

Assignments

Assignment 01: Create a web server and an Amazon RDS DB instance.

Install an Apache web server with PHP and create a MySQL database. The web server runs on an Amazon EC2 instance using Amazon Linux, and the MySQL database is a MySQL DB instance. Both the Amazon EC2 instance and the DB instance run in a virtual private cloud (VPC) based on the Amazon VPC service.

1. Launch an EC2 instance
 2. Create a DB instance
 3. Install a web server on your EC2 instance
1. Create an Amazon EC2 instance in the public subnet of your VPC:
- Sign into the AWS Management Console and open the Amazon EC2 console at <https://console.aws.amazon.com/ec2/> .
 - In the upper-right corner of the AWS Management Console, choose the AWS Region where you want to create the EC2 instance.
 - Choose **EC2 Dashboard**, and then choose **Launch instance**, as shown following
 - Make sure you have opted into the new launch experience
 - Under Name and tags, for Name, enter ec2-instance-web-server.
 - Under **Application and OS Images (Amazon Machine Image)**, choose **Amazon Linux**, and then choose the **Amazon Linux 2 AMI**. Keep the defaults for the other choices.
 - Under **Instance type**, choose **t2.micro**.
 - Under **Key pair (login)**, choose a **Key pair name** to use an existing key pair. To create a new key pair for the Amazon EC2 instance, choose **Create new key pair** and then use the **Create key pair** window to create it.
 - Under **Network settings**, set these values and keep the other values as their defaults:
For **Allow SSH traffic from**, choose the source of SSH connections to the EC2 instance.
Turn on **Allow HTTPS traffic from the internet**.
Turn on **Allow HTTP traffic from the internet**.
 - Leave the default values for the remaining sections.
 - Review a summary of your instance configuration in the **Summary** panel, and when you're ready, choose **Launch instance**.
 - On the **Launch Status** page, shown following, note the identifier for your new EC2 instance.
 - Choose **View all instances** to find your instance.
 - Wait until **Instance state** for your instance is **Running** before continuing.

New EC2 Experience X

EC2 Dashboard
EC2 Global View
Events
Tags
Limits

Instances New
Instance Types
Launch Templates
Spot Requests
Savings Plans
Reserved Instances New
Dedicated Hosts
Scheduled Instances
Capacity Reservations

Images New
AMIs New
AMI Catalog

Elastic Block Store

Feedback Looking for language selection? Find it in the new Unified Settings [?]

Type here to search

© 2022, Amazon Internet Services Private Ltd. or its affiliates. Privacy Terms Cookie preferences

11:25 AM 11/8/2022

Instances (1) Info

Find instance by attribute or tag (case-sensitive)

Name	Instance ID	Instance state	Instance type	Status check	Alarm status	Availability Zone	Public IPv4 DNS
ec2-instance-...	i-0b8db356d57cc4829	Pending	t2.micro	-	No alarms	us-east-1d	ec2-54-167-156-

Select an instance

Manage Approvals - Barclays No X Launch an EC2 instance - Amazon X Instances | EC2 Management Con X +

https://us-east-1.console.aws.amazon.com/ec2/home?region=us-east-1#Instances:instanceState=running

aws Services Search [Alt+S] [?] N. Virginia Nikki Saraswat

New EC2 Experience X
Tell us what you think

EC2 Dashboard
EC2 Global View
Events
Tags
Limits

Instances New
Instance Types
Launch Templates
Spot Requests
Savings Plans
Reserved Instances New
Dedicated Hosts
Scheduled Instances
Capacity Reservations

Images New
AMIs New
AMI Catalog

Elastic Block Store

Feedback Looking for language selection? Find it in the new Unified Settings [?]

Type here to search

© 2022, Amazon Internet Services Private Ltd. or its affiliates. Privacy Terms Cookie preferences

11:29 AM 11/8/2022

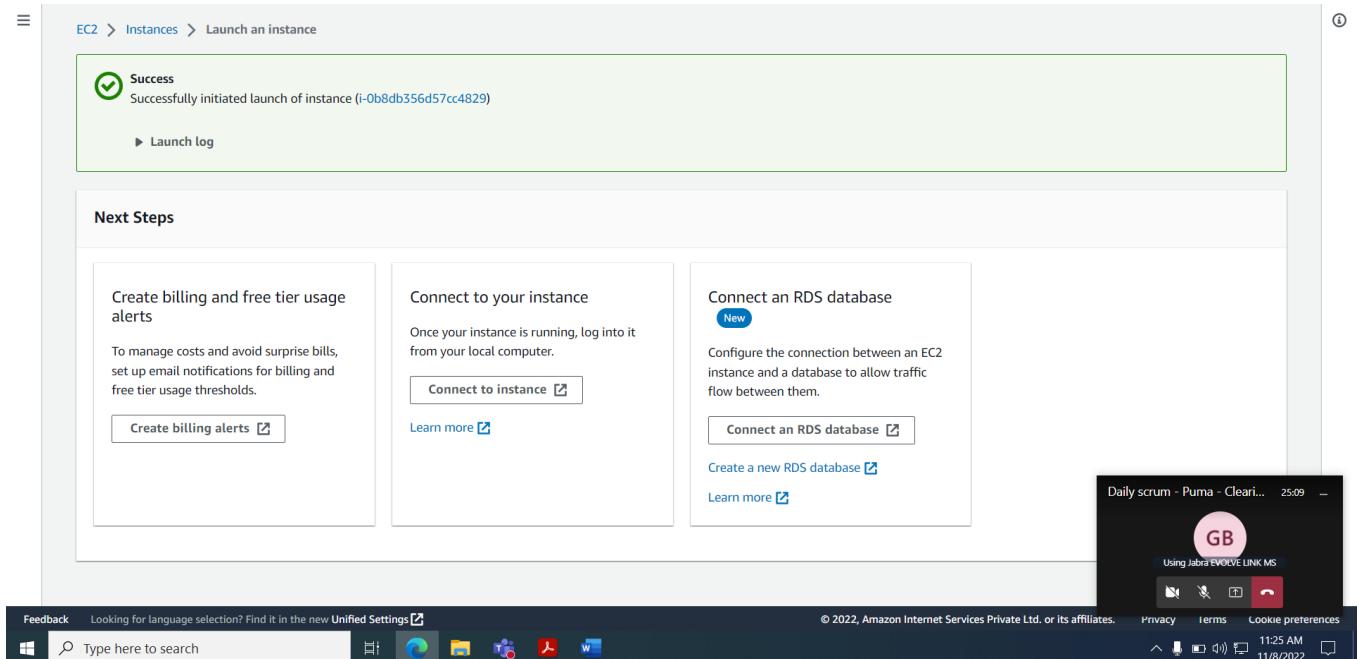
Instances (1) Info

Find instance by attribute or tag (case-sensitive)

Instance state = running X Clear filters

Name	Instance ID	Instance state	Instance type	Status check	Alarm status	Availability Zone	Public IPv4 DNS
ec2-instance-...	i-0b8db356d57cc4829	Running	t2.micro	2/2 checks passed	No alarms	us-east-1d	ec2-54-167-156-

Select an instance



2. Create an Amazon RDS for MySQL DB instance that maintains the data used by a web application.
 - In the upper-right corner of the AWS Management Console, check the AWS Region. It should be the same as the one where you created your EC2 instance
 - In the navigation pane, choose **Databases**.
 - Choose **Create database**.
 - On the **Create database** page, shown following, make sure that the **Standard create** option is chosen, and then choose **MySQL**.
 - In the **Templates** section, choose **Free tier**.
 - In the **Availability and durability** section, keep the defaults.
 - In the **Settings** section, set these values:
DB instance identifier – db-instance
Master username – admin
Auto generate a password – Leave the option turned off.
Master password – *****
Confirm password – *****
 - In the **Instance configuration** section, set these values: **Burstable classes (includes t classes)db.t3.micro**
 - In the **Storage** section, keep the defaults.
 - In the **Connectivity** section, set these values and keep the other values as their defaults:
For Compute resource, choose **Connect to an EC2 compute resource**.
For EC2 instance, choose the EC2 instance you created previously, such as **ec2-instance-web-server**.

- In the **Database authentication** section, make sure **Password authentication** is selected.
- Open the **Additional configuration** section, and enter **sample** for **Initial database name**. Keep the default settings for the other options.
- To create your MySQL DB instance, choose **Create database**. Your new DB instance appears in the **Databases** list with the status **Creating**.
- Wait for the **Status** of your new DB instance to show as **Available**. Then choose the DB instance name to show its details.
- In the **Connectivity & security** section, view the **Endpoint** and **Port** of the DB instance.

The screenshot shows the AWS RDS Management Console interface for creating a new database. The URL in the browser is <https://us-east-1.console.aws.amazon.com/rds/home?region=us-east-1#launch-dbinstance:gdb=false;isHermesCreate=true;s3-import=false>.

We listened to your feedback!

Now, create a database with a single click using our pre-built configurations! Or choose your own configurations.

Create database

Choose a database creation method

- Standard create: You set all of the configuration options, including ones for availability, security, backups, and maintenance.
- Easy create: Use recommended best-practice configurations. Some configuration options can be changed after the database is created.

Engine options

Engine type: [Info](#)

- Amazon Aurora
- MySQL
- MariaDB
- PostgreSQL
- Oracle
- Microsoft SQL Server

PostgreSQL

PostgreSQL is a powerful, open-source object-relational database system with a strong reputation of reliability, stability, and correctness.

- High reliability and stability in a variety of workloads.
- Advanced features to perform in high-volume environments.
- Vibrant open-source community that releases new features multiple times per year.
- Supports multiple extensions that add even more functionality to the database.
- Supports General Purpose, Memory Optimized, and Burstable Performance instance classes.
- The most Oracle-compatible open-source database.

Feedback: Looking for language selection? Find it in the new [Unified Settings](#).

© 2022, Amazon Internet Services Private Ltd. or its affiliates. Privacy Terms Cookie preferences

11:31 AM 11/8/2022

Screenshot of the AWS RDS Management Console showing the 'Create a DB instance' wizard.

The left pane shows the configuration steps:

- Additional configuration**: Database options, encryption turned on, backup turned on, backtrack turned off, maintenance, CloudWatch Logs, delete protection turned off.
- Database options**: Initial database name (sample), DB parameter group (default:mysql8.0), Option group (default:mysql-8-0).
- Backup**: Enable automated backups (checked). A note states: "Please note that automated backups are currently supported for InnoDB storage engine only. If you are using MyISAM, refer to details here." Backup retention period (7 days).

The right pane displays information about PostgreSQL:

- PostgreSQL is a powerful, open-source object-relational database system with a strong reputation of reliability, stability, and correctness.
- High reliability and stability in a variety of workloads.
- Advanced features to perform in high-volume environments.
- Vibrant open-source community that releases new features multiple times per year.
- Supports multiple extensions that add even more functionality to the database.
- Supports General Purpose, Memory Optimized, and Burstable Performance instance classes.
- The most Oracle-compatible open-source database.

Feedback: Looking for language selection? Find it in the new Unified Settings.

© 2022, Amazon Internet Services Private Ltd. or its affiliates. Privacy Terms Cookie preferences

Screenshot of the AWS RDS Management Console showing the 'Create a DB instance' wizard.

The left pane shows the configuration steps:

- Maintenance window**: Select the period you want pending modifications or maintenance applied to the database by Amazon RDS.
 - Choose a window (radio button)
 - No preference (radio button)
- Deletion protection**: Enable deletion protection (checkbox). Note: Protects the database from being deleted accidentally. While this option is enabled, you can't delete the database.
- Estimated monthly costs**: The Amazon RDS Free Tier is available to you for 12 months. Each calendar month, the free tier will allow you to use the Amazon RDS resources listed below for free:
 - 750 hrs of Amazon RDS in a Single-AZ db.t2.micro, db.t3.micro or db.t4g.micro Instance.
 - 20 GB of General Purpose Storage (SSD).
 - 20 GB for automated backup storage and any user-initiated DB Snapshots.

Learn more about AWS Free Tier.

When your free usage expires or if your application use exceeds the free usage tiers, you simply pay standard, pay-as-you-go service rates as described in the [Amazon RDS Pricing page](#).

The right pane displays information about PostgreSQL:

- PostgreSQL is a powerful, open-source object-relational database system with a strong reputation of reliability, stability, and correctness.
- High reliability and stability in a variety of workloads.
- Advanced features to perform in high-volume environments.
- Vibrant open-source community that releases new features multiple times per year.
- Supports multiple extensions that add even more functionality to the database.
- Supports General Purpose, Memory Optimized, and Burstable Performance instance classes.
- The most Oracle-compatible open-source database.

Feedback: Looking for language selection? Find it in the new Unified Settings.

© 2022, Amazon Internet Services Private Ltd. or its affiliates. Privacy Terms Cookie preferences

3. Install a web server on your EC2 instance

[Install an Apache web server with PHP and MariaDB](#)

- Connect to the EC2 instance that you created earlier by following the steps in Connect to your Linux instance.
- Get the latest bug fixes and security updates by updating the software on your EC2 instance. To do this, use the following command.
`sudo yum update -y`
- After the updates complete, install the PHP software using the `amazon-linux-extras` install command. This command installs multiple software packages and related dependencies at the same time.
`sudo amazon-linux-extras install php8.0 mariadb10.5`
- Install the Apache web server.
`sudo yum install -y httpd`
- Start the web server with the command shown following.
`sudo systemctl start httpd`
- Configure the web server to start with each system boot using the `systemctl` command.
`sudo systemctl enable httpd`

To set file permissions for the Apache web server:

- Add the `ec2-user` user to the `apache` group.
`sudo usermod -a -G apache ec2-user`
- Log out to refresh your permissions and include the new `apache` group.
Exit
- Change the group ownership of the `/var/www` directory and its contents to the `apache` group.
`sudo chown -R ec2-user:apache /var/www`
- Recursively change the permissions for files in the `/var/www` directory and its subdirectories to add group write permissions.
`find /var/www -type f -exec sudo chmod 0664 {} \;`

Connect your Apache web server to your DB instance:

- While still connected to your EC2 instance, change the directory to `/var/www` and create a new subdirectory named `inc`.
- Create a new file in the `inc` directory named `dbinfo.inc`, and then edit the file by calling `nano` (or the editor of your choice).
- Add the following contents to the `dbinfo.inc` file. Here, `db_instance_endpoint` is your DB instance endpoint, without the port, and `master password` is the master password for your DB instance.
`db_instance_endpoint = "dbinstanceendpoint";
master_password = "password";`
- Save and close the `dbinfo.inc` file.
- Change the directory to `/var/www/html`.
- Create a new file in the `html` directory named `SamplePage.php`, and then edit the file by calling `nano` (or the editor of your choice).
- Add the following contents to the `SamplePage.php` file:
`<?php
$connection = new mysqli("dbinstanceendpoint", "username", "password");
if ($connection->connect_error) {
 die("Connection failed: " . $connection->connect_error);
}
echo "Connected successfully";
</?php`
- Save and close the `SamplePage.php` file.

- Verify that your web server successfully connects to your DB instance by opening a web browser and browsing to `http://EC2 instance endpoint/SamplePage.php`

The screenshot shows the AWS EC2 Instance Connect interface. At the top, there are several tabs: "Manage Approvals - Barclays No", "Install a web server on your EC2", "RDS Management Console", "Connect to instance | EC2 Manager", and "EC2 Instance Connect". Below the tabs, the AWS navigation bar shows "Services" and "Search". The main content area is titled "Connect to instance" and displays the following information:

- Instance ID:** i-0b8db356d57cc4829 (ec2-instance-web-server)
- Public IP address:** 54.167.156.74
- User name:** ec2-user
- Note:** In most cases, the guessed user name is correct. However, read your AMI usage instructions to check if the AMI owner has changed the default AMI user name.

At the bottom right are "Cancel" and "Connect" buttons.

The screenshot shows the AWS EC2 Instance Connect interface with a progress bar indicating the transfer of a file named "file.tar.gz". The progress bar shows the file size as 1.2 GB and the transfer time as 00:00:00. The progress is at 0% completion.

Below the progress bar, the file transfer details are listed:

File	Size	Transfer Time
file.tar.gz	1.2 GB	00:00:00

At the bottom left, it says "i-0b8db356d57cc4829 (ec2-instance-web-server)" and "PublicIPs: 54.167.156.74 PrivateIPs: 172.31.83.251".

The screenshot shows the AWS EC2 Instance Connect interface with a progress bar indicating the transfer of a file named "file.tar.gz". The progress bar shows the file size as 1.2 GB and the transfer time as 00:00:00. The progress is at 0% completion.

Below the progress bar, the file transfer details are listed:

File	Size	Transfer Time
file.tar.gz	1.2 GB	00:00:00

At the bottom left, it says "i-0b8db356d57cc4829 (ec2-instance-web-server)" and "PublicIPs: 54.167.156.74 PrivateIPs: 172.31.83.251".

```
Installing : mod_http2-1.15.19-1.amzn2.0.1.x86_64          8/9
Installing : httpd-2.4.54-1.amzn2.x86_64                  9/9
Verifying  : apr-util-1.6.1-5.amzn2.0.2.x86_64           1/9
Verifying  : apr-util-bdb-1.6.1-5.amzn2.x86_64           2/9
Verifying  : httpd-tools-2.4.54-1.amzn2.x86_64            3/9
Verifying  : mod_http2-1.15.19-1.amzn2.0.1.x86_64          4/9
Verifying  : httpd-2.4.54-1.amzn2.x86_64                  5/9
Verifying  : mailcap-2.1.41-2.amzn2.noarch                6/9
Verifying  : generic-logos-httpd-18.0.0-4.amzn2.noarch     7/9
Verifying  : httpd-filesystem-2.4.54-1.amzn2.noarch        8/9
Verifying  : apr-1.7.0-9.amzn2.x86_64                      9/9

Installed:
  httpd.x86_64 0:2.4.54-1.amzn2

Dependency Installed:
  apr.x86_64 0:1.7.0-9.amzn2      apr-util.x86_64 0:1.6.1-5.amzn2.0.2  apr-util-bdb.x86_64 0:1.6.1-5.amzn2.0.2  generic-logos-httpd.noarch 0:18.0.0-4.amzn2
  httpd-filesystem.noarch 0:2.4.54-1.amzn2  httpd-tools.x86_64 0:2.4.54-1.amzn2  mailcap.noarch 0:2.1.41-2.amzn2  mod_http2.x86_64 0:1.15.19-1.amzn2.0.1

Complete!
[ec2-user@ip-172-31-83-251 ~]$ sudo systemctl start httpd
[ec2-user@ip-172-31-83-251 ~]$ sudo systemctl start httpd
[ec2-user@ip-172-31-83-251 ~]$ sudo systemctl enable httpd
Created symlink from /etc/systemd/system/multi-user.target.wants/httpd.service to /usr/lib/systemd/system/httpd.service.
[ec2-user@ip-172-31-83-251 ~]$ sudo usermod -a -G apache ec2-user
[ec2-user@ip-172-31-83-251 ~]$ exit
logout
```

i-0b8db356d57cc4829 (ec2-instance-web-server)

PublicIPs: 54.167.156.74 PrivateIPs: 172.31.83.251

Feedback Looking for language selection? Find it in the new Unified Settings

© 2022, Amazon Internet Services Private Ltd. or its affiliates.

Privacy

Terms

Cookie preferences

11:40 AM 11/8/2022

```
Last login: Tue Nov  8 06:06:37 2022 from ec2-18-206-107-28.compute-1.amazonaws.com
[ec2-user@ip-172-31-83-251 ~]$ groups
ec2-user adm wheel apache systemd-journal
[ec2-user@ip-172-31-83-251 ~]$ sudo chown -R ec2-user:apache /var/www
[ec2-user@ip-172-31-83-251 ~]$
```

i-0b8db356d57cc4829 (ec2-instance-web-server)

PublicIPs: 54.167.156.74 PrivateIPs: 172.31.83.251

Feedback Looking for language selection? Find it in the new Unified Settings

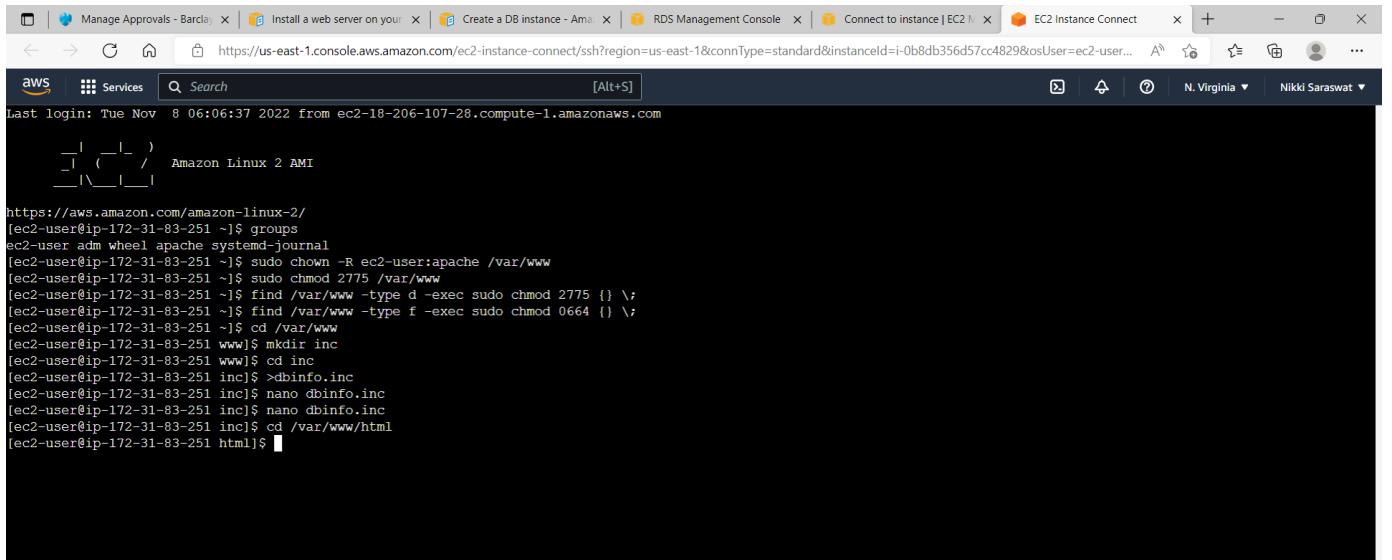
© 2022, Amazon Internet Services Private Ltd. or its affiliates.

Privacy

Terms

Cookie preferences

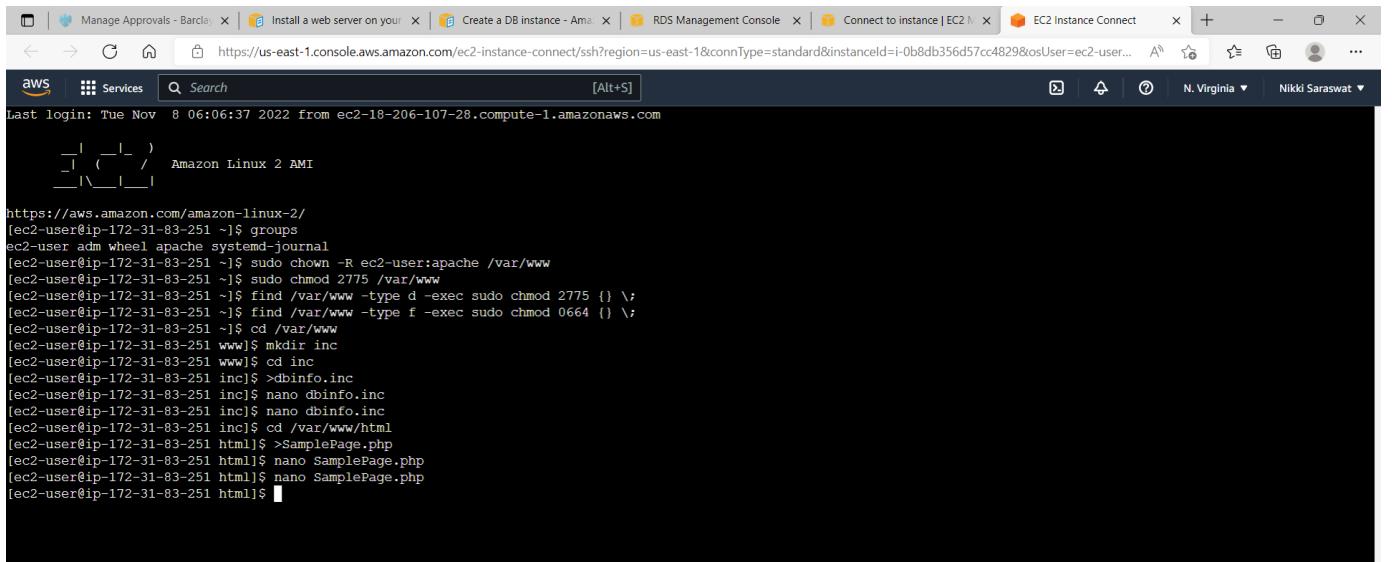
11:42 AM 11/8/2022



```
Last login: Tue Nov  8 06:06:37 2022 from ec2-18-206-107-28.compute-1.amazonaws.com
[ec2-user@ip-172-31-83-251 ~]$ groups
ec2-user adm wheel apache systemd-journal
[ec2-user@ip-172-31-83-251 ~]$ sudo chown -R ec2-user:apache /var/www
[ec2-user@ip-172-31-83-251 ~]$ sudo chmod 2775 /var/www
[ec2-user@ip-172-31-83-251 ~]$ find /var/www -type d -exec sudo chmod 2775 {} \;
[ec2-user@ip-172-31-83-251 ~]$ find /var/www -type f -exec sudo chmod 0664 {} \;
[ec2-user@ip-172-31-83-251 ~]$ cd /var/www
[ec2-user@ip-172-31-83-251 www]$ mkdir inc
[ec2-user@ip-172-31-83-251 www]$ cd inc
[ec2-user@ip-172-31-83-251 inc]$ >dbinfo.inc
[ec2-user@ip-172-31-83-251 inc]$ nano dbinfo.inc
[ec2-user@ip-172-31-83-251 inc]$ nano dbinfo.inc
[ec2-user@ip-172-31-83-251 inc]$ cd /var/www/html
[ec2-user@ip-172-31-83-251 html]$ nano SamplePage.php
[ec2-user@ip-172-31-83-251 html]$ nano SamplePage.php
[ec2-user@ip-172-31-83-251 html]$ nano SamplePage.php
```

i-0b8db356d57cc4829 (ec2-instance-web-server)

PublicIPs: 54.167.156.74 PrivateIPs: 172.31.83.251

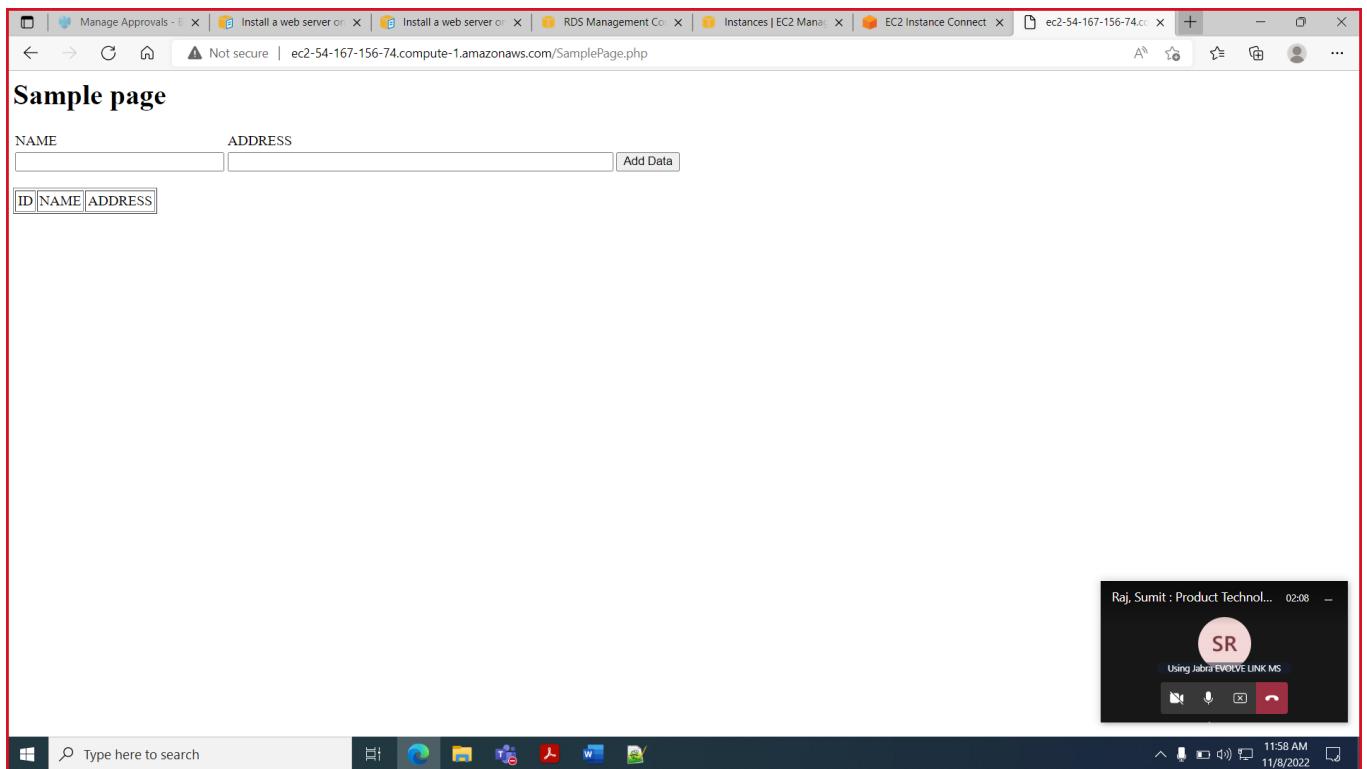


```
Last login: Tue Nov  8 06:06:37 2022 from ec2-18-206-107-28.compute-1.amazonaws.com
[ec2-user@ip-172-31-83-251 ~]$ groups
ec2-user adm wheel apache systemd-journal
[ec2-user@ip-172-31-83-251 ~]$ sudo chown -R ec2-user:apache /var/www
[ec2-user@ip-172-31-83-251 ~]$ sudo chmod 2775 /var/www
[ec2-user@ip-172-31-83-251 ~]$ find /var/www -type d -exec sudo chmod 2775 {} \;
[ec2-user@ip-172-31-83-251 ~]$ find /var/www -type f -exec sudo chmod 0664 {} \;
[ec2-user@ip-172-31-83-251 ~]$ cd /var/www
[ec2-user@ip-172-31-83-251 www]$ mkdir inc
[ec2-user@ip-172-31-83-251 www]$ cd inc
[ec2-user@ip-172-31-83-251 inc]$ >dbinfo.inc
[ec2-user@ip-172-31-83-251 inc]$ nano dbinfo.inc
[ec2-user@ip-172-31-83-251 inc]$ nano dbinfo.inc
[ec2-user@ip-172-31-83-251 inc]$ cd /var/www/html
[ec2-user@ip-172-31-83-251 html]$ nano SamplePage.php
[ec2-user@ip-172-31-83-251 html]$ nano SamplePage.php
[ec2-user@ip-172-31-83-251 html]$ nano SamplePage.php
```

i-0b8db356d57cc4829 (ec2-instance-web-server)

PublicIPs: 54.167.156.74 PrivateIPs: 172.31.83.251





Assignment 02: (Amplify, Lambda & DynamoDB)

In this tutorial, you will create a simple web application. You will first build a static web app that renders "HelloWorld." Then you will learn how to add functionality to the web app so the text that displays is based on a custom input you provide.

Steps:

1. Create Web App: Deploy static resources for your web application using the AWS Amplify Console.
2. Build Serverless Function: Build a serverless function using AWS Lambda.
3. Link Serverless Function to Web App: Deploy your serverless function with API Gateway.
4. Create Data Table: Persist data in an Amazon DynamoDB table.
5. Add Interactivity to Web App: Modify your web app to invoke your API.

1. Create Web App with Amplify Console

- Open your favorite text editor on your computer. Create a new file and paste the following HTML in it:

```
<!DOCTYPE html>
<html>
```

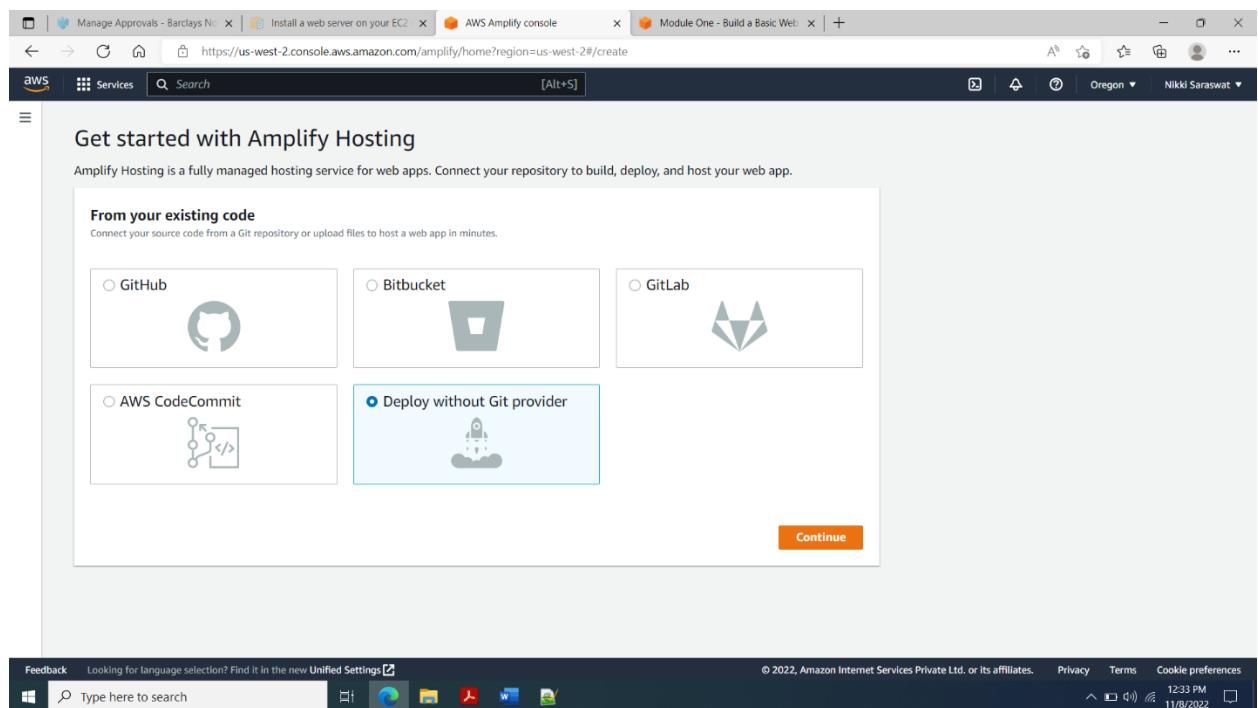
```

<head>
  <meta charset="UTF-8">
  <title>Hello World</title>
</head>

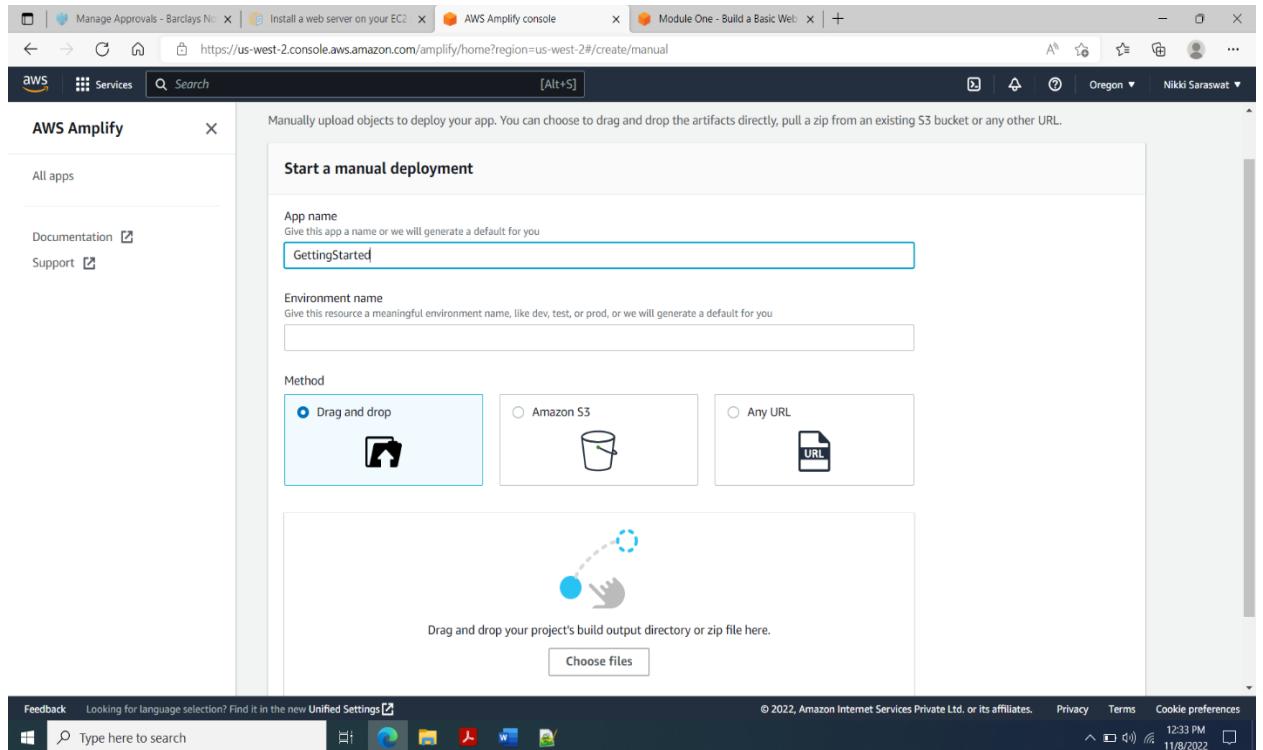
<body>
  Hello World
</body>
</html>

```

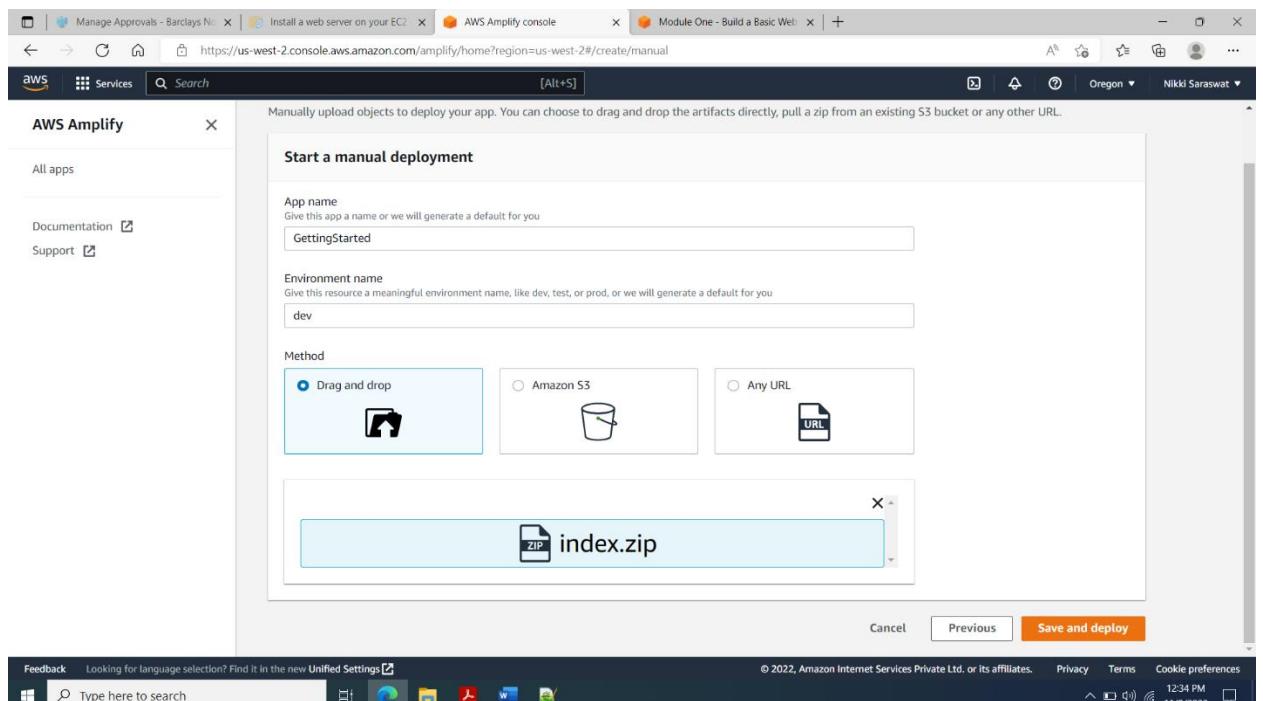
- Save the file as index.html.
- In a new browser window, log into the Amplify console. Note: We will be using the Oregon (us-west-2) Region for this tutorial.]



- In the Get Started section, under Host your web app, choose the orange Get started button.
- Select Deploy without Git provider.
- Choose the Continue button.
- In the App name field, enter GettingStarted.



- For Environment name, enter dev.
- Select the Drag and drop method.
- Choose the Choose files button.
- Select the ZIP file you created in Step 3.



- Choose the Save and deploy button.
- After a few seconds, you should see the message Deployment successfully completed.

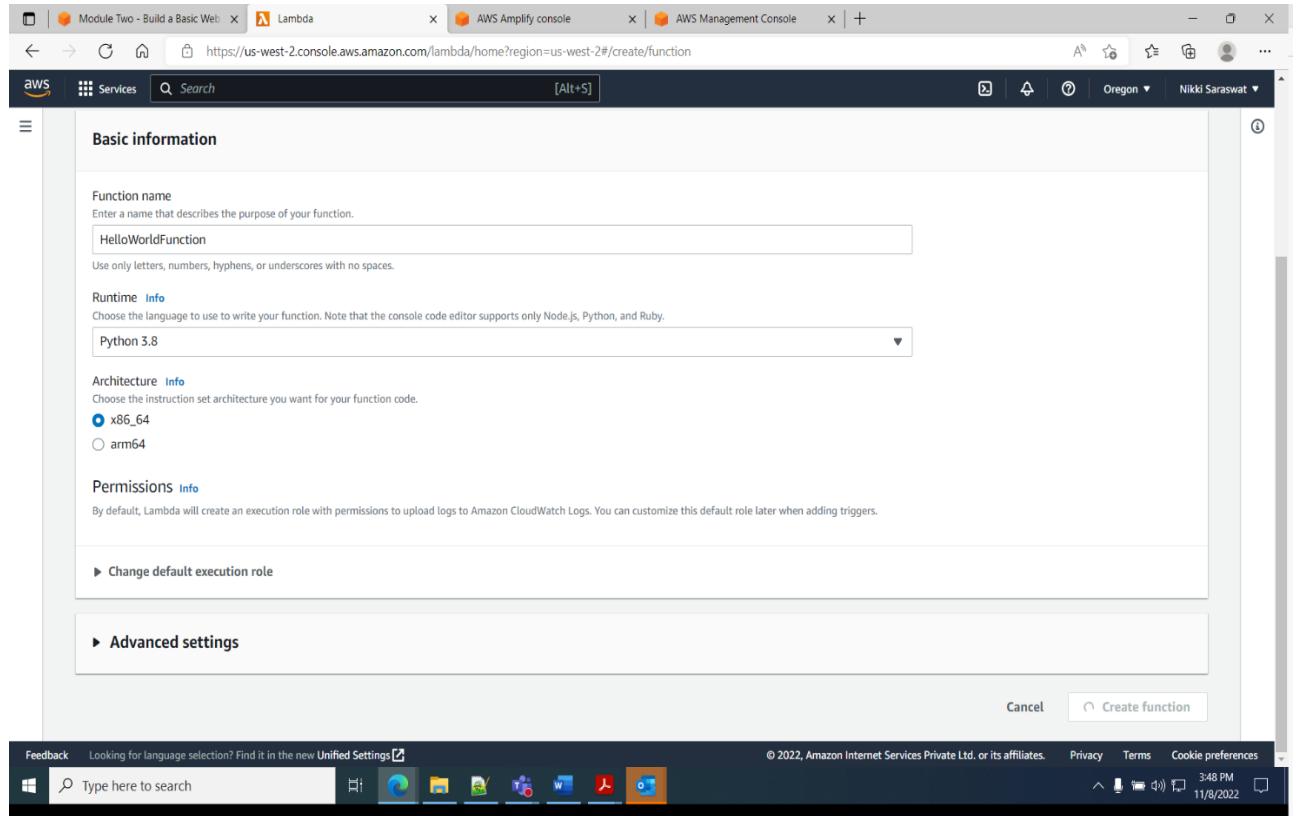
The screenshot shows the AWS Amplify console interface. On the left, a sidebar titled "AWS Amplify" lists "All apps" (GettingStarted) and "App settings" (General, Amplify Studio settings, Domain management, Notifications, Access control, Monitoring, Rewrites and redirects, Custom headers). Under "Domain management", it says "amplifyapp.com" is available. The main content area shows a table with one row:

Domain	Status
amplifyapp.com	Available

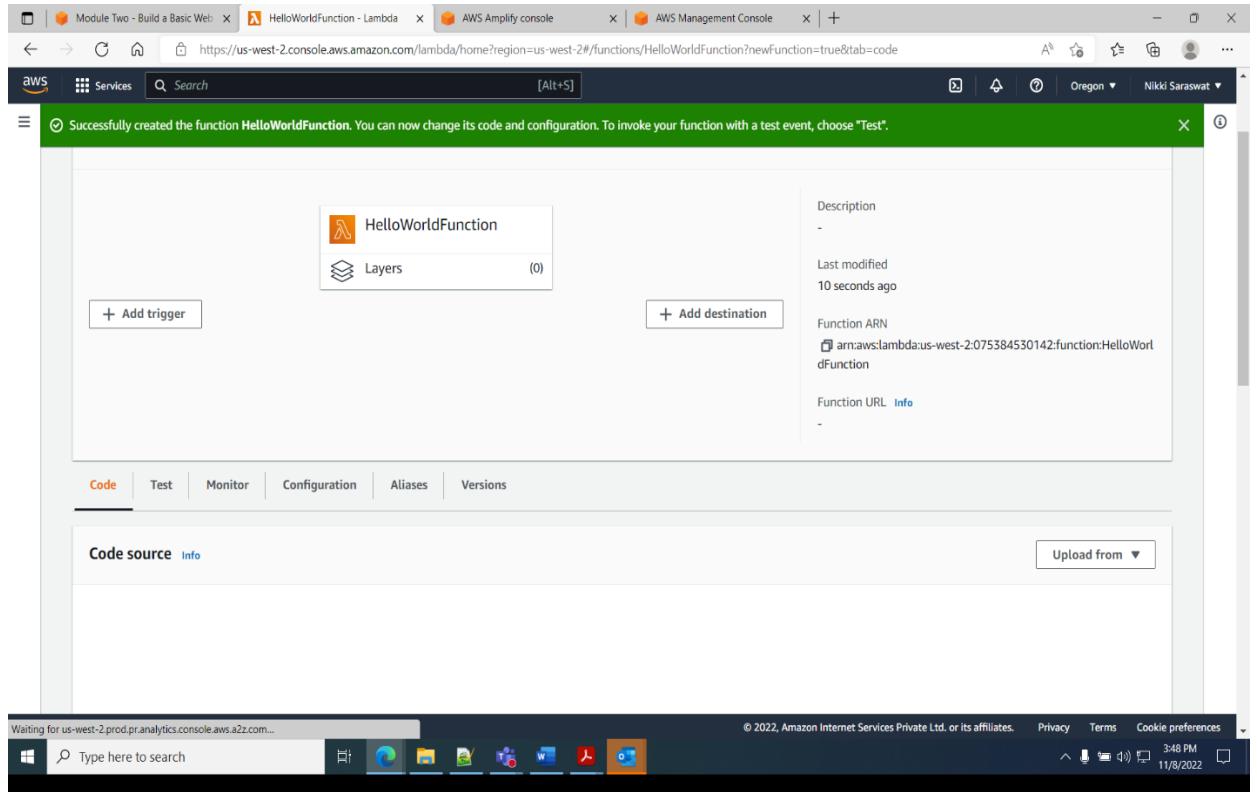
Below the table, there's a URL entry: "https://dev.d17olobxt0dvs0.amplifyapp.com" with "Branch" set to "dev" and "Redirects to" set to "-". At the bottom of the page, there's a Windows taskbar with icons for File Explorer, Task View, Edge browser, File Manager, and others.

2. Build a Serverless Function

- In a new browser tab, log in to the AWS Lambda console.
- Make sure you create your function in the same Region in which you created the web app in the previous module. You can see this at the very top of the page, next to your account name.
- Choose the orange Create function button.
- Under Function name, enter HelloWorldFunction.



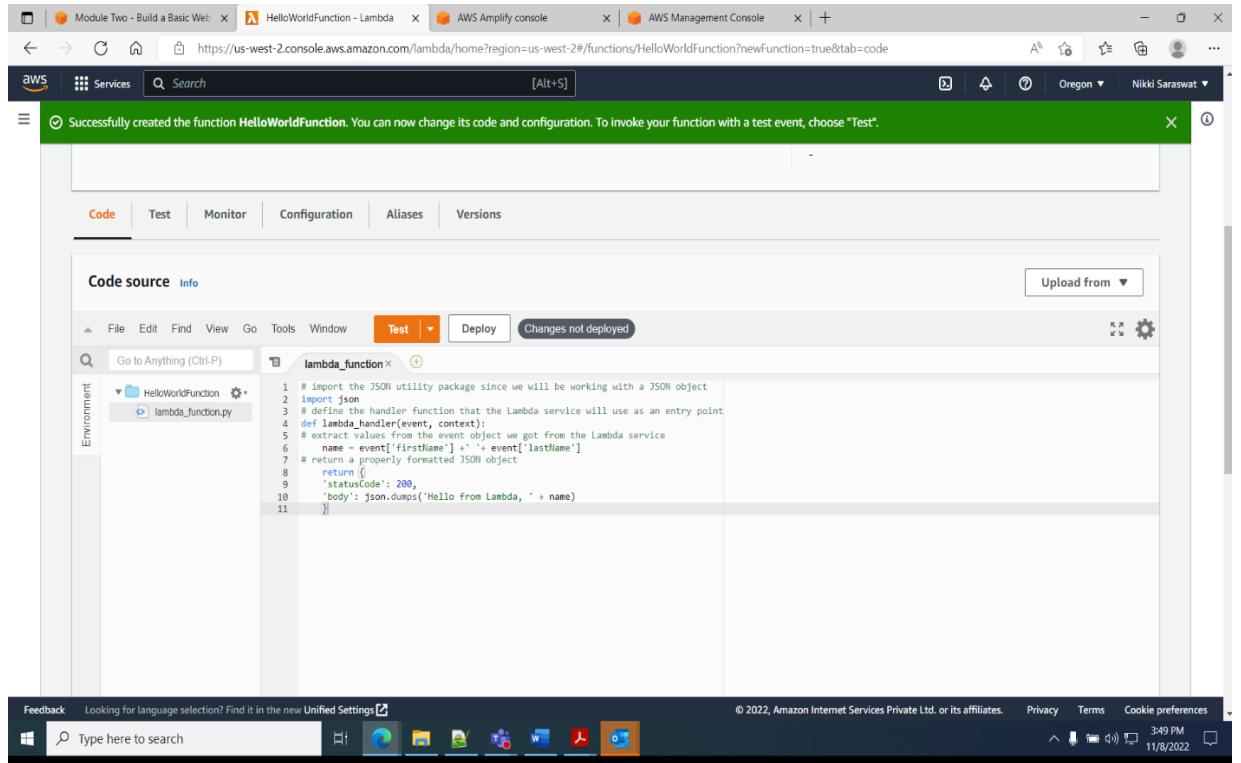
- Select Python 3.8 from the runtime dropdown and leave the rest of the defaults unchanged.
- Choose the orange Create function button.
- You should see a green message box at the top of your screen with the following message "Successfully created the function HelloWorldFunction."



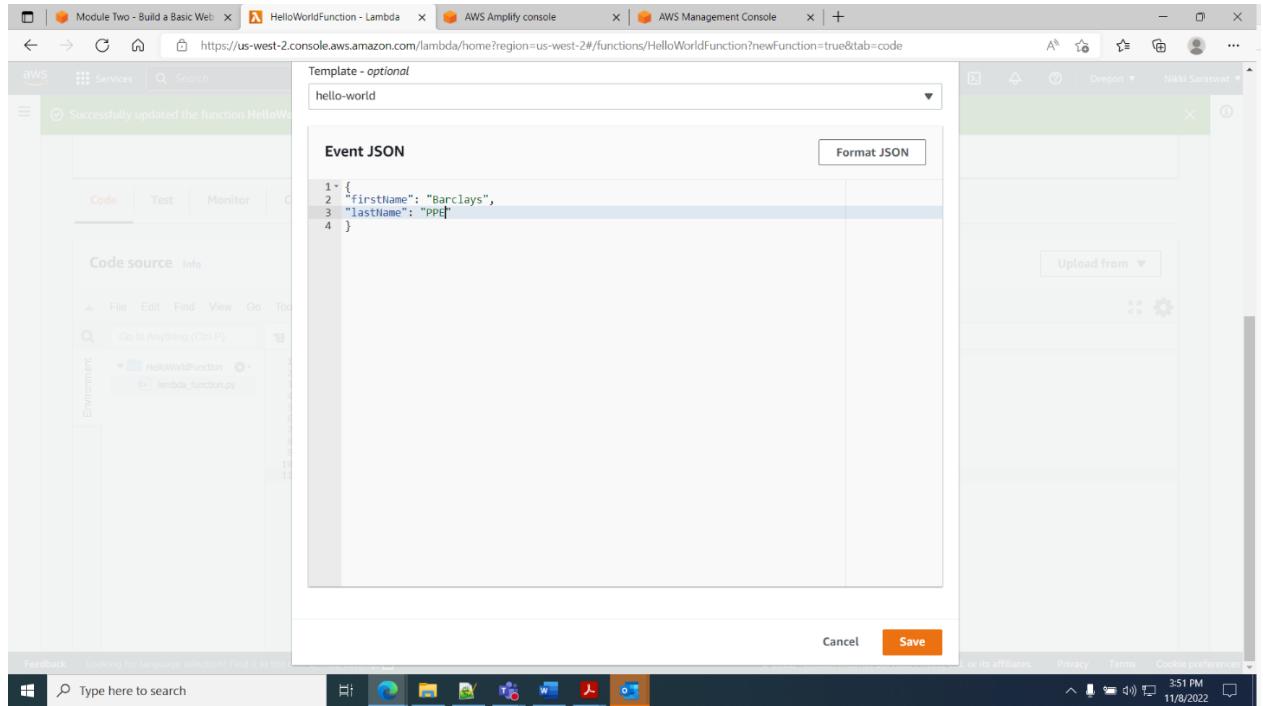
- Under Code source, replace the code in lambda_function.py with the following:

```
# import the JSON utility package since we will be working with a JSON object
import json

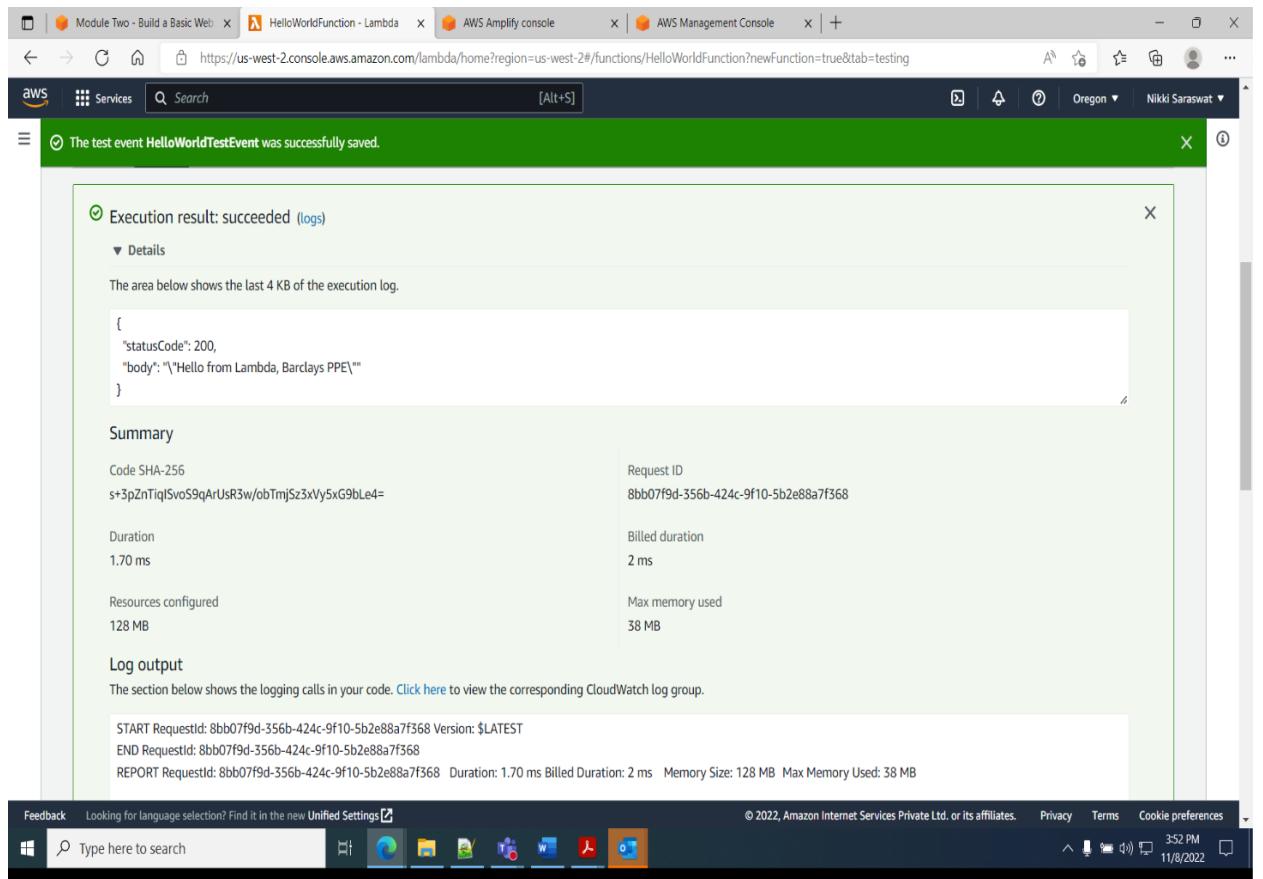
# define the handler function that the Lambda service will use as an entry point
def lambda_handler(event, context):
    # extract values from the event object we got from the Lambda service
    name = event['firstName'] + ' ' + event['lastName']
    # return a properly formatted JSON object
    return {
        'statusCode': 200,
        'body': json.dumps('Hello from Lambda, ' + name)
    }
```



- Save by going to the file menu and selecting Save to save the changes.
- Choose Deploy to deploy the changes.
- Let's test our new function. Choose the orange Test button to create a test event by selecting Configure test event.
- Under Event name, enter HelloWorldTestEvent.
- Copy and paste the following JSON object to replace the default one:

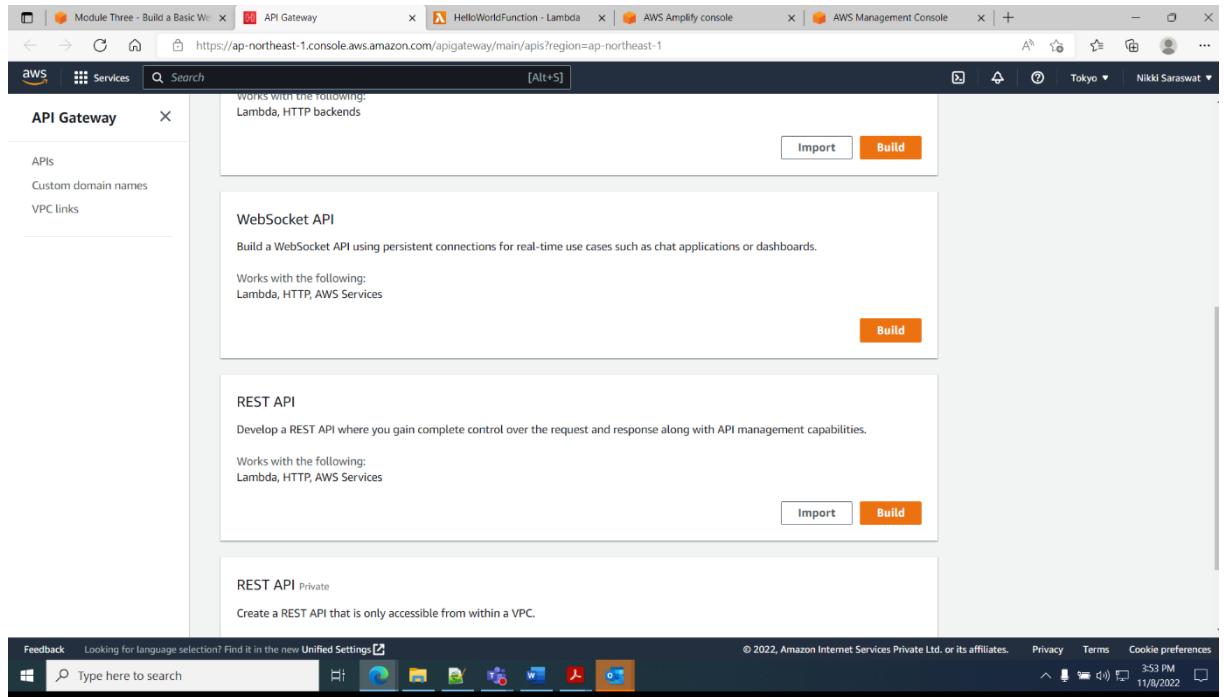


- Choose the orange Create button at the bottom of the page.

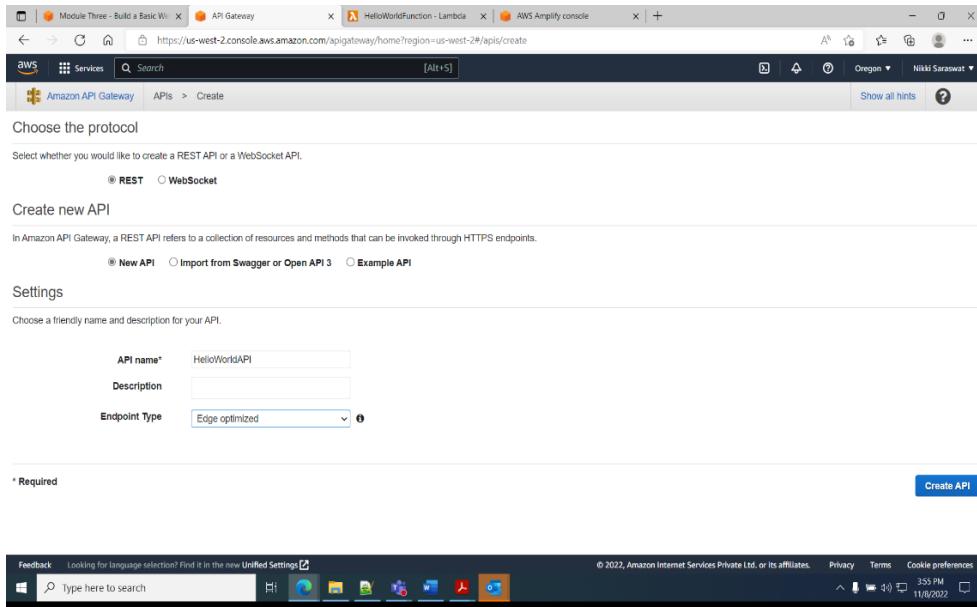


3. Link a Serverless Function to a Web App

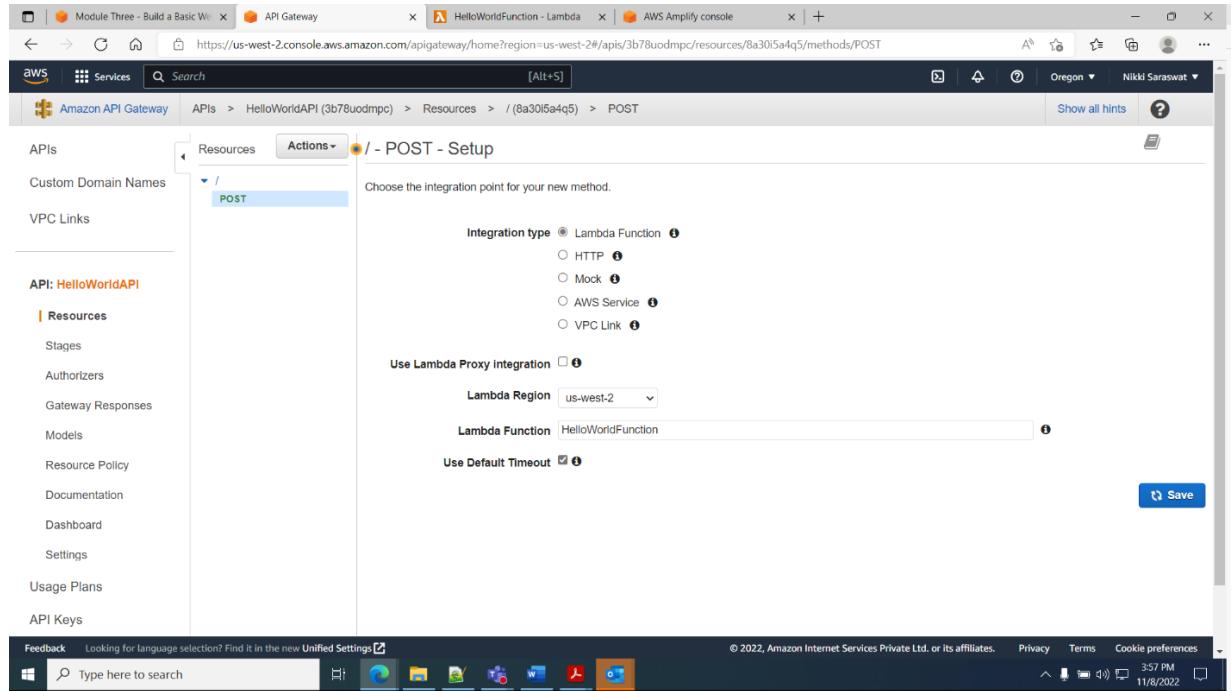
- Log in to the API Gateway console.



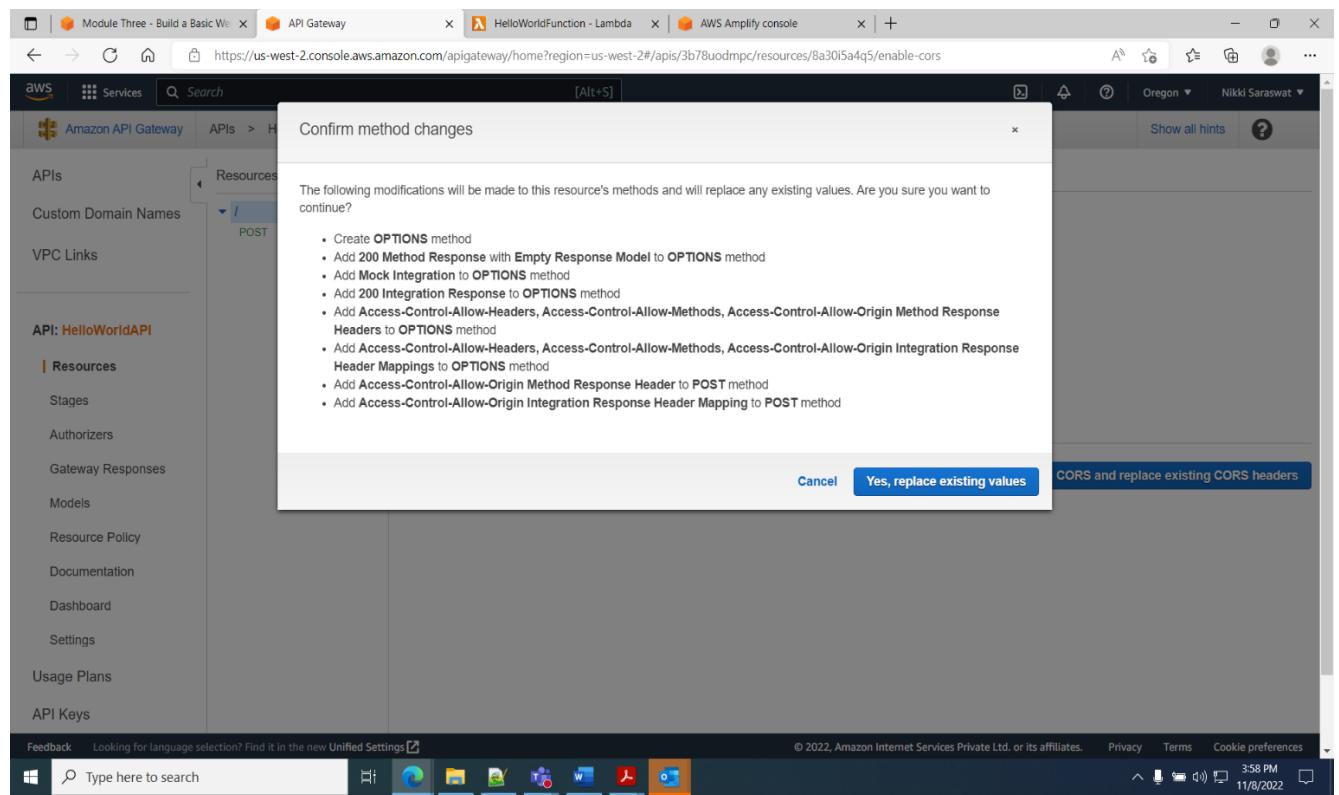
- In the Choose an API type section, find the REST API card and choose the Build button on the card.
- Under Choose the protocol, select REST.
- Under Create new API, select New API.
- In the API name field, enter HelloWorldAPI.
- Select Edge optimized from the Endpoint Type dropdown. (Note: Edge-optimized endpoints are best for geographically distributed clients. This makes them a good choice for public services being accessed from the internet. Regional endpoints are typically used for APIs that are accessed primarily from within the same AWS Region.)
- Choose the blue Create API button. Your settings should look like the accompanying screenshot.



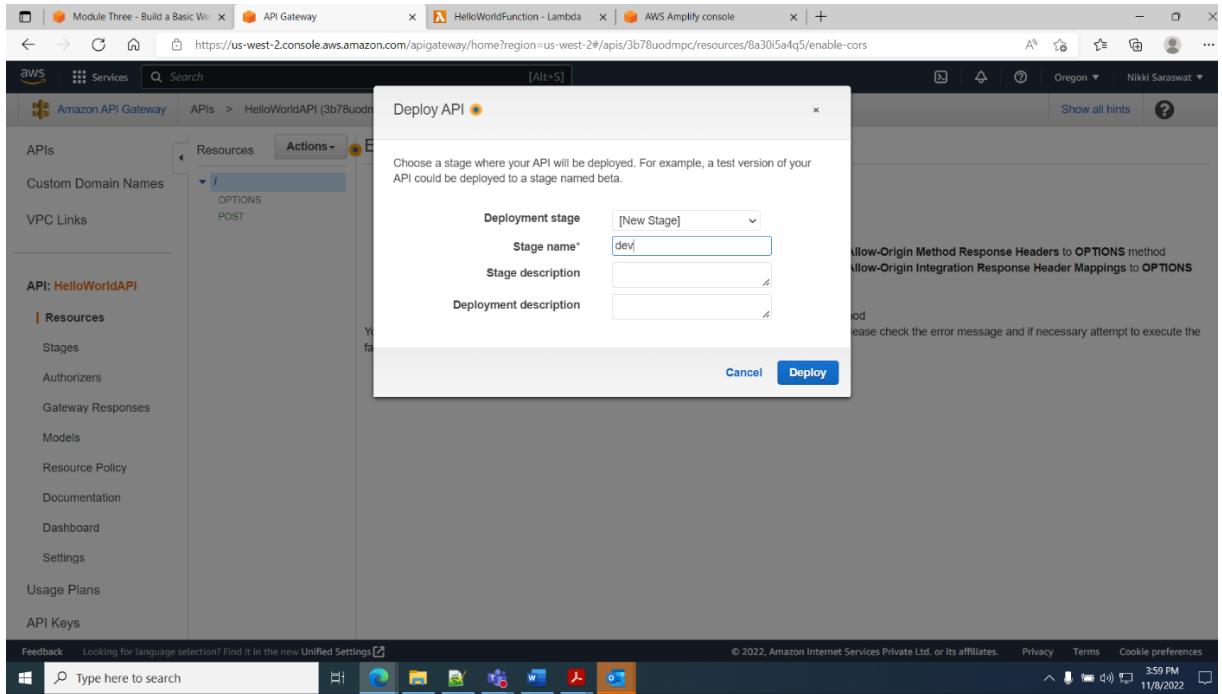
- In the left navigation pane, select Resources under API: HelloWorldAPI.
- Ensure the "/" resource is selected.
- From the Actions dropdown menu, select Create Method.
- Select POST from the new dropdown that appears, then select the checkmark.
- Select Lambda Function for the Integration type.
- Select the Lambda Region you used when making the function (or else you will see a warning box reading "You do not have any Lambda Functions in...").
- Enter HelloWorldFunction in the Lambda Function field.
- Choose the blue Save button.
- You should see a message letting you know you are giving the API you are creating permission to call your Lambda function. Choose the OK button.
- With the newly created POST method selected, select Enable CORS from the Action dropdown menu.
- Leave the POST checkbox selected and choose the blue Enable CORS and replace existing CORS headers button.



- You should see a message asking you to confirm method changes. Choose the blue Yes, replace existing values button.



- In the Actions dropdown list, select Deploy API.
- Select [New Stage] in the Deployment stage dropdown list.
- Enter dev for the Stage Name.
- Choose Deploy.
- Copy and save the URL next to Invoke URL (you will need it in module five).



- In the left navigation pane, select Resources.
- The methods for our API will now be listed on the right. Choose POST.
- Choose the small blue lightning bolt.
- Paste the following into the Request Body field:
- Choose the blue Test button.
- On the right side, you should see a response with Code 200.
- Great! We have built and tested an API that calls our Lambda function.

The screenshot shows the AWS API Gateway console with a POST method configuration. The request body is set to:

```

1- {
2-   "firstName": "Barclays",
3-   "lastName": "Posting"
4- }

```

The screenshot shows the AWS API Gateway console under the 'Method Execution' tab for a POST request. The response body is:

```
{"statusCode": 200, "body": "Hello from Lambda, Barclays Posting"}
```

4. Create a Data Table

- Log in to the Amazon DynamoDB console.
- Make sure you create your table in the same Region in which you created the web app in the previous module. You can see this at the very top of the page, next to your account name.

- Choose the orange Create table button.
- Under Table name, enter HelloWorldDatabase.
- In the Partition key field, enter ID. The partition key is part of the table's primary key.
- Leave the rest of the default values unchanged and choose the orange Create table button.
- In the list of tables, select the table name, HelloWorldDatabase.
- In the General information section, show Additional info by selecting the down arrow.
- Copy the Amazon Resource Name (ARN). You will need it later in this module.

Table details

DynamoDB is a schemaless database that requires only a table name and a primary key when you create the table.

Table name
This will be used to identify your table.

Between 3 and 255 characters, containing only letters, numbers, underscores (_), hyphens (-), and periods (.)

Partition key
The partition key is part of the table's primary key. It is a hash value that is used to retrieve items from your table and allocate data across hosts for scalability and availability.

Sort key - optional
You can use a sort key as the second part of a table's primary key. The sort key allows you to sort or search among all items sharing the same partition key.

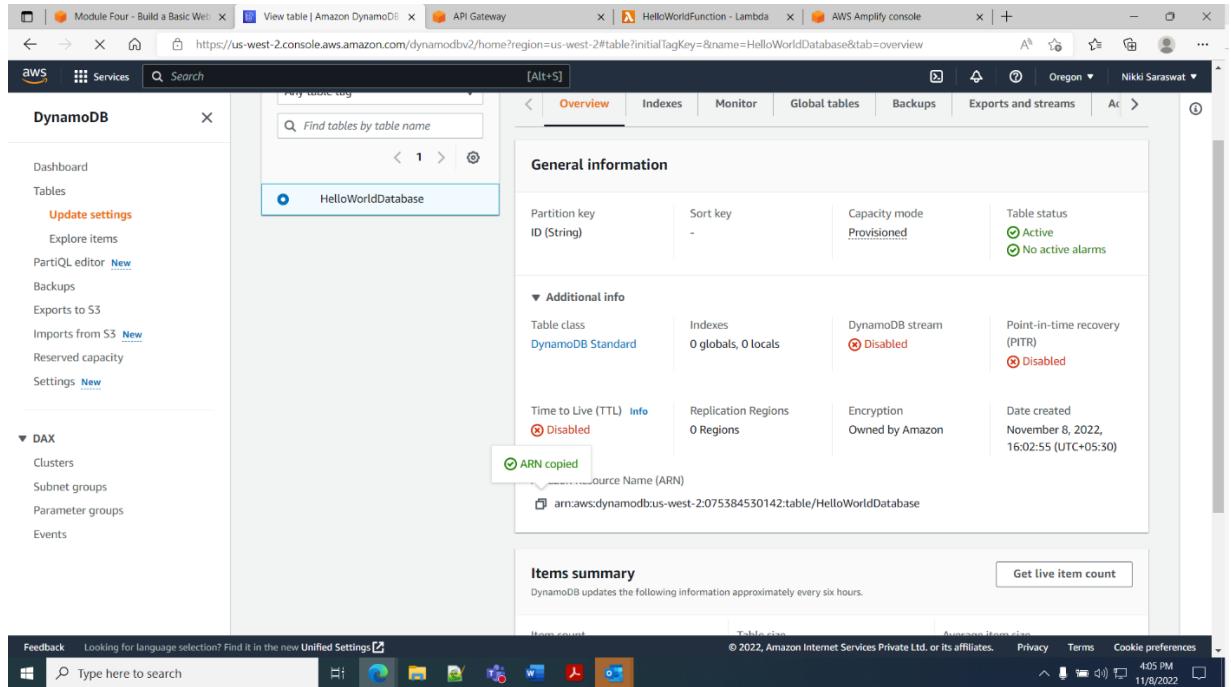
Table settings

Capacity mode	Provisioned	Yes
Read capacity	5 RCU	Yes
Write capacity	5 WCU	Yes
Auto scaling	On	Yes
Local secondary indexes	-	No
Global secondary indexes	-	Yes
Encryption key management	Owned by Amazon DynamoDB	Yes
Table class	DynamoDB Standard	Yes

Tags
Tags are pairs of keys and optional values, that you can assign to AWS resources. You can use tags to control access to your resources or track your AWS spending.

No tags are associated with the resource.

You can add 50 more tags.



- Now that we have a table, let's edit our Lambda function to be able to write data to it. In a new browser window, open the AWS Lambda console.
- Select the function we created in module two (if you have been using our examples, it will be called HelloWorldFunction). If you don't see it, check the Region dropdown in the upper right next to your name to ensure you're in the same Region you created the function in.
- We'll be adding permissions to our function so it can use the DynamoDB service, and we'll be using AWS Identity and Access Management (IAM) to do so.
- Select the Configuration tab and select Permissions from the right side menu.
- In the Execution role box, under Role name, choose the link. A new browser tab will open.
- In the Permissions policies box, open the Add permissions dropdown and select Create inline policy.
- Select the JSON tab.
- Paste the following policy in the text area, taking care to replace your table's ARN in the Resource field in line 15:
- This policy will allow our Lambda function to read, edit, or delete items, but restrict it to only be able to do so in the table we created.
- Choose the blue Review Policy button.
- Next to Name, enter HelloWorldDynamoPolicy.
- Choose the blue Create Policy button.
- You can now close this browser tab and go back to the tab for your Lambda function.

```

1 {
2   "Version": "2012-10-17",
3   "Statement": [
4     {
5       "Sid": "VisualEditor0",
6       "Effect": "Allow",
7       "Action": [
8         "dynamodb:PutItem",
9         "dynamodb>DeleteItem",
10        "dynamodb:GetItem",
11        "dynamodb:Scan",
12        "dynamodb:Query",
13        "dynamodb:UpdateItem"
14      ],
15      "Resource": "arn:aws:dynamodb:us-west-2:075384530142:table/HelloWorldDatabase"
16    }
17  }
18 }

```

Character count: 283 of 10,240.
The current character count includes character for all inline policies in the role: HelloWorldFunction-role-yuox9295.

Feedback: Looking for language selection? Find it in the new Unified Settings. © 2022, Amazon Internet Services Private Ltd. or its affiliates. Privacy Terms Cookie preferences

Type here to search

Character count: 283 of 10,240.
The current character count includes character for all inline policies in the role: HelloWorldFunction-role-yuox9295.

Feedback: Looking for language selection? Find it in the new Unified Settings. © 2022, Amazon Internet Services Private Ltd. or its affiliates. Privacy Terms Cookie preferences

Type here to search

- Select the Code tab and select your function from the navigation pane on the left side of the code editor.
- Replace the code for your function with the following:
- Choose the Deploy button at the top of the code editor.

The screenshot shows the AWS Lambda console interface. The main area displays the code source for the function `HelloWorldFunction`. The code is written in Python and uses the AWS SDK (boto3) to interact with a DynamoDB table named `HelloWorldDatabase`. It defines a handler function `lambda_handler` that takes an event object and returns a response. The response includes the name from the event and a timestamp. The AWS Lambda service is used as an entry point for the function.

```
# Import the json utility package since we will be working with a JSON object
import json
# Import the AWS SDK (for Python the package name is boto3)
import boto3
# Import time
import time
# Import datetime, strftime
from time import strftime, strptime
# Create a DynamoDB object using the AWS SDK
dynamodb = boto3.resource('dynamodb')
# Use the DynamoDB object to select our table
table = dynamodb.Table('HelloWorldDatabase')
# Store the current date and time in a variable
now = strftime("%X %d %B %Y %H:%M:%S +0000", gettime())
# Define the handler function that the Lambda service will use as an entry point
def lambda_handler(event, context):
    # Extract values from the event object we got from the Lambda service and store in a variable
    name = event['firstname'] + ' ' + event['lastname']
    # Write the name and time to the DynamoDB table using the object we instantiated and save response in a variable
    response = table.put_item(
        Item={
            'ID': name,
            'LatestGreetingTime': now
        }
    )
    # Return a properly formatted JSON object
    return {
        "statusCode": 200,
        "body": json.dumps('Hello from Lambda, ' + name)
    }

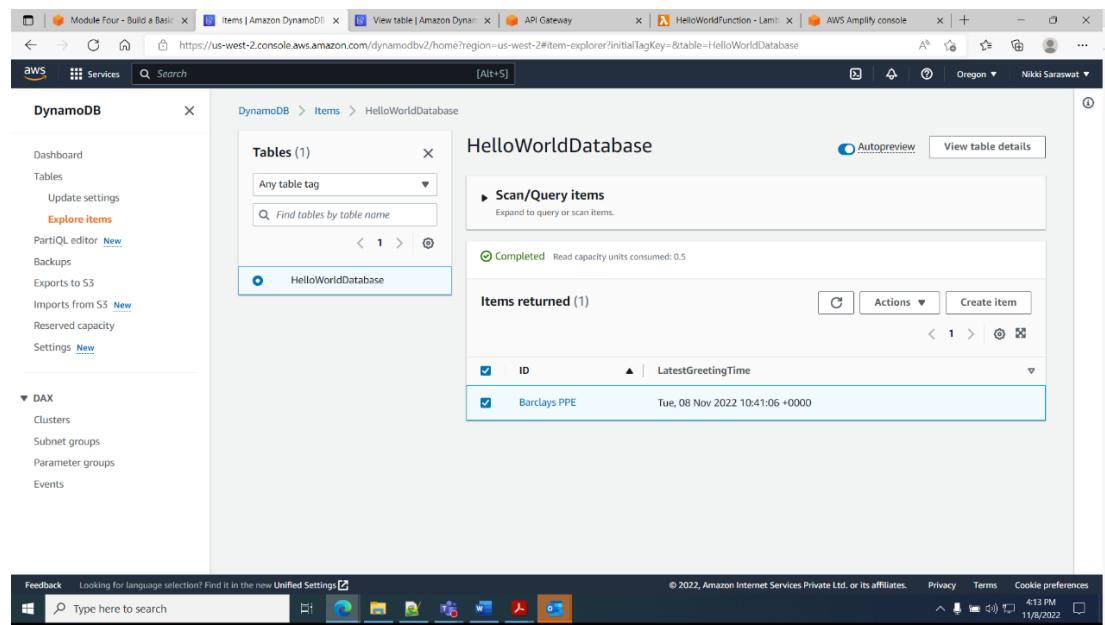
```

The screenshot shows the AWS Lambda console interface after a successful deployment. A green banner at the top indicates that the function `HelloWorldFunction` has been successfully updated. Below this, the code source is shown again, and the `Test` tab is selected. The `Execution result` section displays the test event name `HelloWorldTestEvent` and the response. The response is a JSON object with a status code of 200 and a body containing the message "Hello from Lambda, Barclays PPE". The `Function logs` section shows the request ID, start and end request IDs, duration, and memory usage for the test execution.

Successfully updated the function `HelloWorldFunction`.

```
{
  "statusCode": 200,
  "body": "Hello from Lambda, Barclays PPE"
}
```

Request ID: b050bbf0-71f2-45e6-86eb-9bdc1418d56 Duration: 217.05 ms Billed Duration: 218 ms Memory Size: 128 MB Max Memory Used: 64 MB Init Duration: 338.99 ms



5. Add Interactivity to Your Web App

- Open the index.html file you created in module one.
- Replace the existing code with the following:

```
<!DOCTYPE html>
<html>
<head>
<meta charset="UTF-8">
<title>Hello World</title>
<!-- Add some CSS to change client UI -->
<style>
body {
    background-color: #232F3E;
}
label, button {
    color: #FF9900;
    font-family: Arial, Helvetica, sans-serif;
    font-size: 20px;
    margin-left: 40px;
}
input {
    color: #232F3E;
    font-family: Arial, Helvetica, sans-serif;
    font-size: 20px;
    margin-left: 20px;
}
</style>
<script>
// define the callAPI function that takes a first name and last name as parameters
var callAPI = (firstName,lastName)=>{
    // instantiate a headers object
    var myHeaders = new Headers();
    // add content type header to object
    myHeaders.append("Content-Type", "application/json");

```

```

// using built in JSON utility package turn object to string and store in a variable
var raw = JSON.stringify({ "firstName":firstName,"lastName":lastName});
// create a JSON object with parameters for API call and store in a variable
var requestOptions = {
  method: 'POST',
  headers: myHeaders,
  body: raw,
  redirect: 'follow'
};
// make API call with parameters and use promises to get response
fetch("YOUR-API-Invoke-URL", requestOptions)
.then(response => response.text())
.then(result => alert(JSON.parse(result).message))
.catch(error => console.log('error', error));
}
</script>
</head>
<body>
<form>
<label>First Name :</label>
<input type="text" id="fName">
<label>Last Name :</label>
<input type="text" id="lName">
<!-- set button onClick method to call function we defined passing input values as parameters -->
<button type="button"
onclick="callAPI(document.getElementById('fName').value,document.getElementById('lName').value)">Call
API</button>
</form>
</body>
</html>

```

- Make sure you add your API Invoke URL on Line 41 (from module three). Note: If you do not have your API's URL, you can get it from the API Gateway console by selecting your API and choosing stages.
- Save the file.
- ZIP (compress) only the HTML file.
- Open the Amplify console.
- Choose the web app created in module one.
- Choose the white Choose files button.
- Select the ZIP file you created in Step 5.
- When the file is uploaded, a deployment process will automatically begin. Once you see a greenbar, your deployment will be complete.

This tab lists all connected branches, select a branch to view build details.

dev

Deployment successfully completed.

Domain: <https://dev.d1c7yshpy5stfn.amplifyapp.com> | Last deployment: 11/8/2022, 4:15:42 PM

Drag and drop your project's build output directory or zip file here to update your app, or, choose another method.

Choose files

All apps > GettingStarted > App settings: Domain management

Domain management

Use your own custom domain with free HTTPS to provide a secure, friendly URL for your app. Register your domain on Amazon Route53 for a one-click setup, or connect any domain registered on a 3rd party provider.

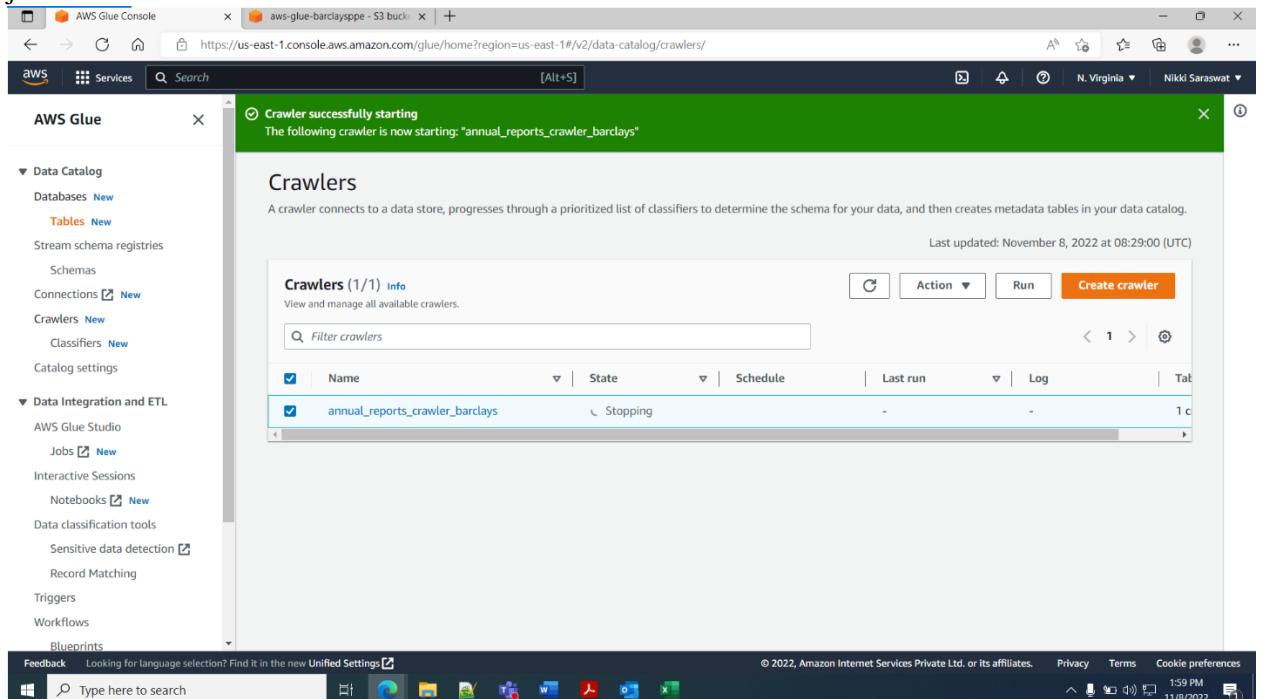
Domain	Status
amplifyapp.com	Available

URL	Branch	Redirects to
https://dev.d1c7yshpy5stfn.amplifyapp.com	dev	-

Assignment 03: GLUE CRAWLER

- Creation of IAM Role for Glue with AWSGlueServiceRole and S3FullAccess Permission

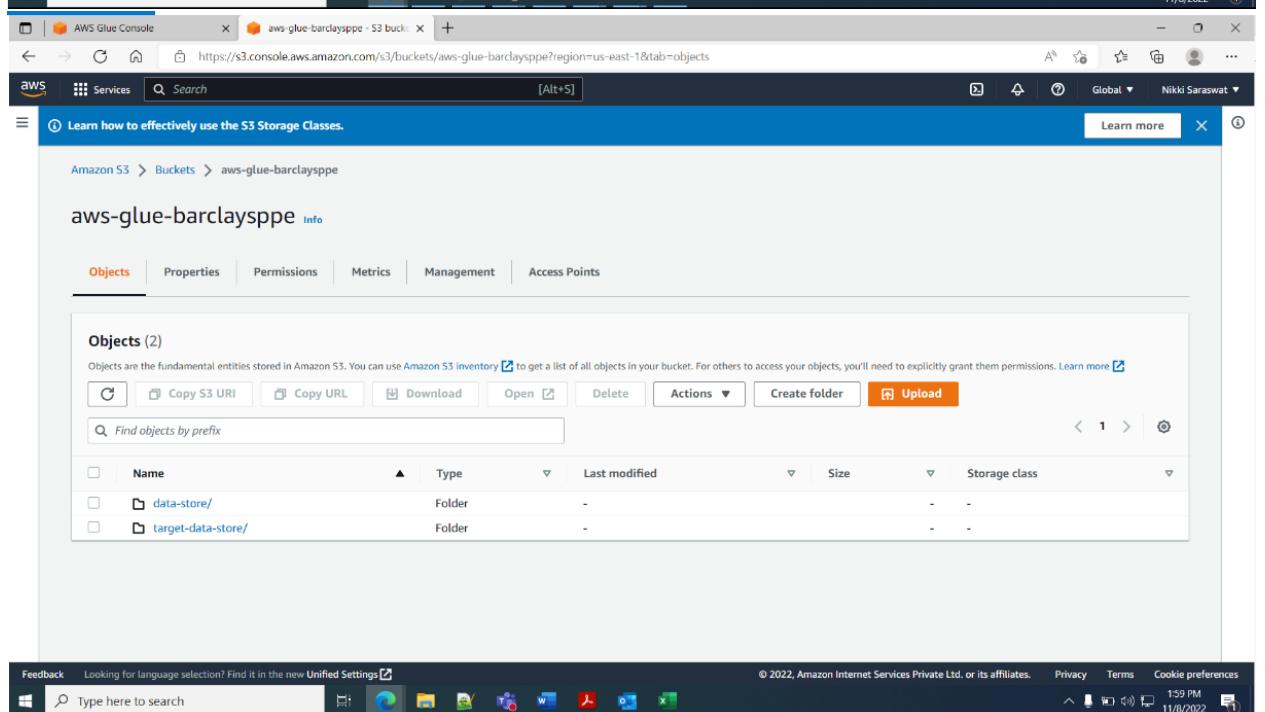
- Now, IN S3 directory structure needs to be created. First directory will be created along with uploaded csv file. Second Structure will be created without any file.
- Now in Glue, database configuration, table information and manage glue studio for declaring jobs needs to be done.



The screenshot shows the AWS Glue console interface. A modal window at the top center displays the message: "Crawler successfully starting" and "The following crawler is now starting: *annual_reports_crawler_barclays*". The main panel is titled "Crawlers" and contains a table with one row:

Name	State	Last run	Log	Tat
annual_reports_crawler_barclays	Stopping	-	-	1 c

On the left sidebar, under "Data Catalog", there are links for Databases, Tables, and Crawlers. Under "Data Integration and ETL", there are links for AWS Glue Studio, Jobs, Notebooks, and Workflows.



The screenshot shows the AWS S3 console interface. The URL is https://s3.console.aws.amazon.com/s3/buckets/aws-glue-barclaysppe?region=us-east-1&tab=objects. The main panel is titled "aws-glue-barclaysppe" and shows the "Objects" tab selected. It displays two objects in the "Objects (2)" section:

Name	Type	Last modified	Size	Storage class
data-store/	Folder	-	-	-
target-data-store/	Folder	-	-	-

At the bottom of the page, there is a feedback bar: "Feedback Looking for language selection? Find it in the new Unified Settings" and a copyright notice: "© 2022, Amazon Internet Services Private Ltd. or its affiliates. Privacy Terms Cookie preferences".

AWS Glue Console - aws-glue-barclayspp - S3 buck | + https://us-east-1.console.aws.amazon.com/glue/home?region=us-east-1#/data-catalog/crawlers/ AWS Services Search [Alt+S] N. Virginia Nikki Saraswat

Crawler successfully starting
The following crawler is now starting: "annual_reports_crawler_barclays"

Crawlers

A crawler connects to a data store, progresses through a prioritized list of classifiers to determine the schema for your data, and then creates metadata tables in your data catalog.

Last updated: November 8, 2022 at 08:29:42 (UTC)

Crawlers (1/1) Info View and manage all available crawlers.

Action Run Create crawler

Filter crawlers

Name	State	Schedule	Last run	Log	Table changes from last r...
annual_reports_crawler_barclays	Ready			Succeeded... View log	1 created

Data Catalog Databases New Tables New Stream schema registries Schemas Connections New Crawlers New Classifiers New Catalog settings Data Integration and ETL AWS Glue Studio Jobs New Interactive Sessions Notebooks New Data classification tools Sensitive data detection Record Matching Triggers Workflows Blueprints Feedback Looking for language selection? Find it in the new Unified Settings © 2022, Amazon Internet Services Private Ltd. or its affiliates. Privacy Terms Cookie preferences Type here to search 1:59 PM 11/8/2022

AWS Glue Studio - aws-glue-barclayspp - S3 buck | + https://us-east-1.console.aws.amazon.com/gluestudio/home?region=us-east-1#/jobs AWS Services Search [Alt+S] N. Virginia Nikki Saraswat

AWS Glue Studio > Jobs

Jobs Info

Create job Info Create

Visual with a source and target Start with a source, ApplyMapping transform, and target.

Visual with a blank canvas Author using an interactive visual interface.

Spark script editor Write or upload your own Spark code.

Python Shell script editor Write or upload your own Python shell script.

Jupyter Notebook Write your own code in a Jupyter Notebook for interactive development.

Source Amazon S3 JSON, CSV, or Parquet files stored in S3. Target Amazon S3 S3 bucket by specifying a bucket path as the data target.

Your jobs (0) Info Actions Run job

Find jobs

Job name	Type	Last modified	AWS Glue version
No jobs			

Feedback Looking for language selection? Find it in the new Unified Settings © 2022, Amazon Internet Services Private Ltd. or its affiliates. Privacy Terms Cookie preferences Type here to search 2:00 PM 11/8/2022

AWS Glue Studio - S3 bucket

https://us-east-1.console.aws.amazon.com/gluestudio/home?region=us-east-1#jobs

N. Virginia Nikki Sarawat

Jobs

Create job

Visual with a source and target

Start with a source, ApplyMapping transform, and target.

Visual with a blank canvas

Author using an interactive visual interface.

Spark script editor

Write or upload your own Spark code.

Python Shell script editor

Write or upload your own Python shell script.

Jupyter Notebook

Write your own code in a Jupyter Notebook for interactive development.

Your jobs (0)

Find jobs

Actions Run job

No jobs

You have not created a job yet.

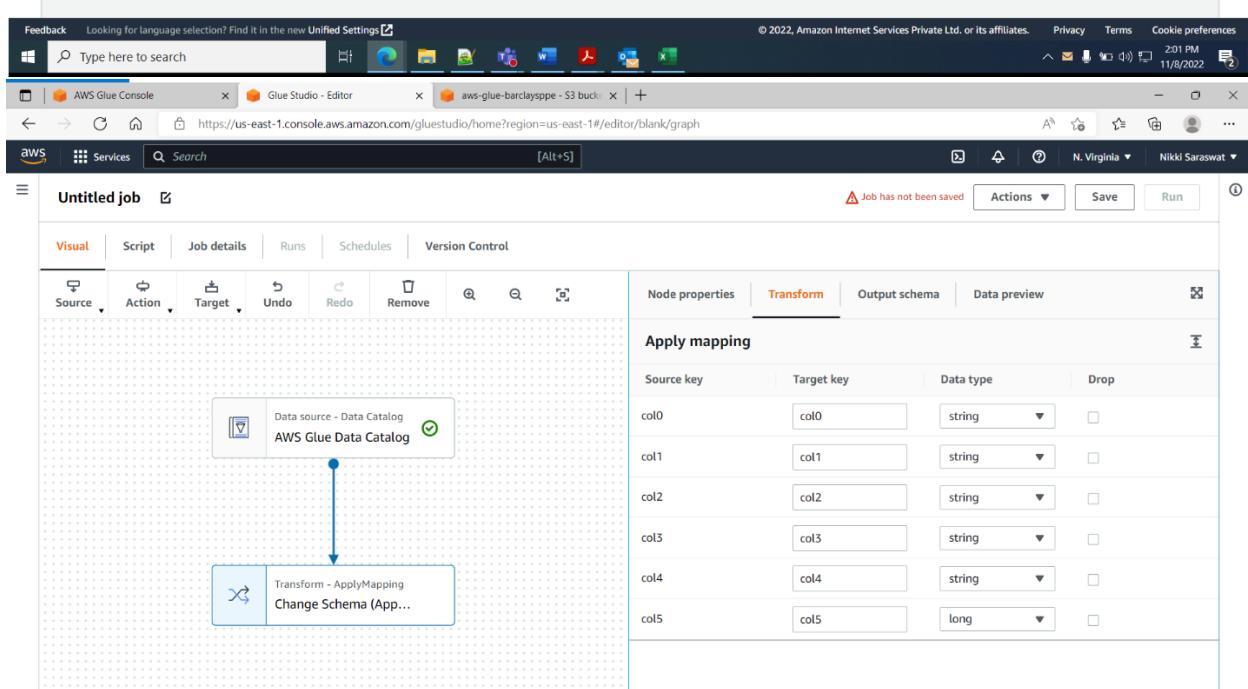
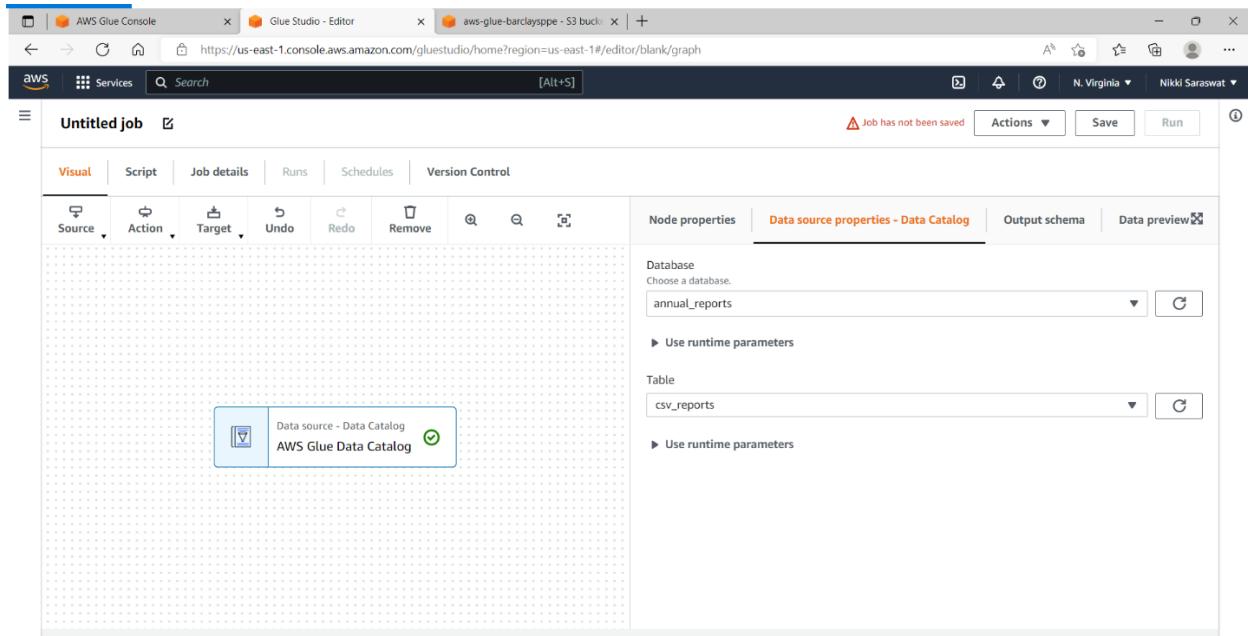
Create job from a blank graph

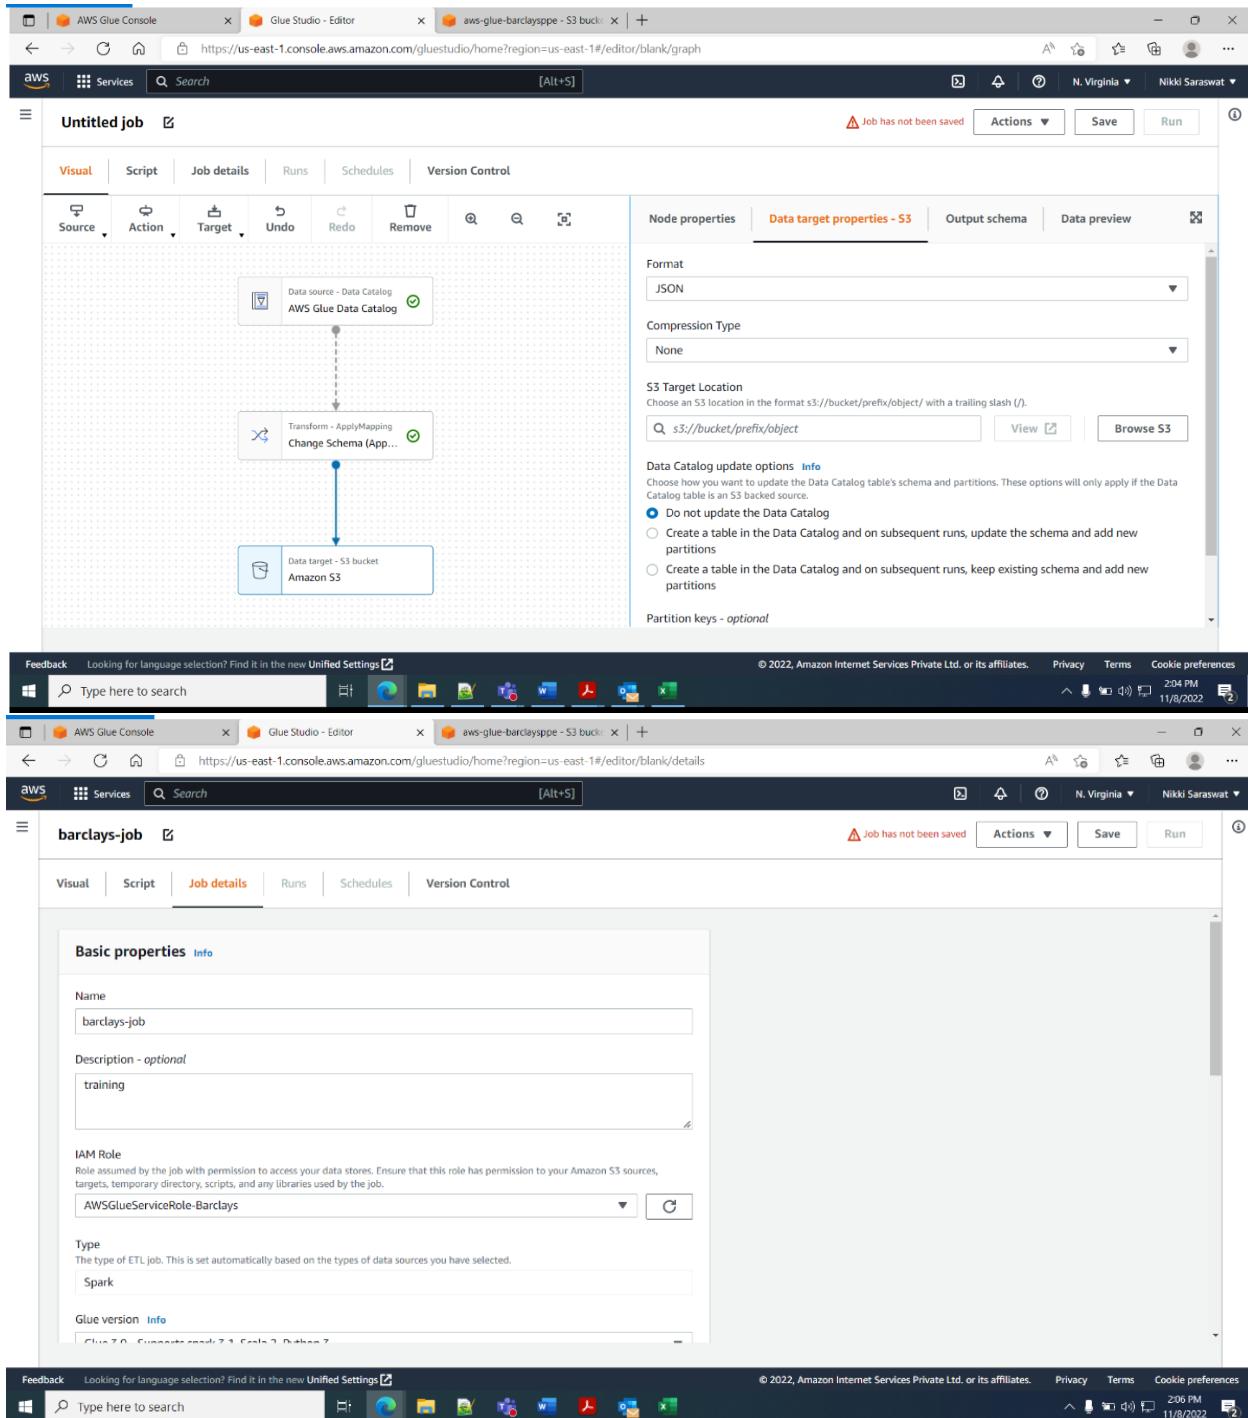
Feedback Looking for language selection? Find it in the new Unified Settings

© 2022, Amazon Internet Services Private Ltd. or its affiliates. Privacy Terms Cookie preferences

2900 PM 11/8/2022

The screenshot shows the AWS Glue Studio interface. At the top, there are three tabs: 'AWS Glue Console', 'AWS Glue Studio', and 'aws-glue-barclaysppe - S3 buck'. The main content area is titled 'Jobs' with a 'Create' button. It displays four options for creating a job: 'Visual with a source and target', 'Visual with a blank canvas' (which is selected), 'Spark script editor', and 'Python Shell script editor'. Below this is a section titled 'Your jobs (0)' with a search bar and a 'Create job from a blank graph' button. The bottom of the screen shows the Windows taskbar with various pinned icons and system status information.





The screenshot shows the AWS Glue Console interface. A green success message at the top states: "Successfully updated job barclays-job. Successfully updated job csv_reports. To run the job choose the Run Job button." Below this, the "Job details" tab is selected. The "Basic properties" section includes fields for Name (barclays-job), Description (optional) (training), IAM Role (AWSGlueServiceRole-Barclays), and Type (Spark). The "Feedback" bar at the bottom suggests looking for language selection in Unified Settings.

barclays-job

Successfully updated job barclays-job. Successfully updated job csv_reports. To run the job choose the Run Job button.

Visual | Script | **Job details** | Runs | Schedules | Version Control

Basic properties

Name: barclays-job

Description - optional: training

IAM Role: AWSGlueServiceRole-Barclays

Type: Spark

Feedback: Looking for language selection? Find it in the new Unified Settings.

The screenshot shows the S3 Management Console. It displays input settings for an S3 Select query. The path is s3://aws-glue-barclayspe/target-data-store/parquet-reports/run-Amaz... and the format is set to Apache Parquet. The output settings section shows a single row with a CSV format. The Feedback bar at the bottom suggests looking for language selection in Unified Settings.

Learn how to effectively use the S3 Storage Classes.

Query with S3 Select

Use Amazon S3 Select to retrieve a subset of data from an object using standard SQL queries. Pricing is based on the size of the input, query results, and data transferred. Learn more or see Amazon S3 pricing.

Input settings

Path: s3://aws-glue-barclayspe/target-data-store/parquet-reports/run-Amaz...
Size: 157.0 B
Format: CSV JSON Apache Parquet
Compression: Amazon S3 Select does not support whole-object compression for Apache Parquet objects.

Output settings

Format: CSV

Feedback: Looking for language selection? Find it in the new Unified Settings.

AWS Glue Console | Glue Studio - Editor | S3 Management Console

https://s3.console.aws.amazon.com/s3/buckets/aws-glue-barclaysppe/object/select?region=us-east-1&prefix=target-data-store/parquet-reports/run-AmazonS3_no... [Alt+S]

aws Services Search [Alt+S] Global Nikki Saraswat

Learn how to effectively use the S3 Storage Classes.

SQL query

Amazon S3 Select supports only the SELECT SQL command. Using the S3 console, you can extract up to 40 MB of records from an object that is up to 128 MB in size. To work with larger files or more records, use the AWS CLI, AWS SDK, or Amazon S3 REST API. For more complex SQL queries, use Amazon Athena.

Add SQL from templates Run SQL query

```
1 /* To create reference point for writing SQL queries, you can display the first 5 records of input data by running the following SQL query: SELECT * FROM s3object s LIMIT 5 */  
2 SELECT * FROM s3object s LIMIT 5
```

Query results

Query results are not available after you choose Close or navigate away. Choose Download results to download a copy of the following query results.

Status

Successfully returned 0 records in 3023 ms

Bytes returned: 0 B

Download results

Close

Feedback Looking for language selection? Find it in the new Unified Settings

Type here to search

© 2022, Amazon Internet Services Private Ltd. or its affiliates. Privacy Terms Cookie preferences 2:11 PM 11/8/2022